

McGill **Artificial Intelligence** Society



Lecture 1: Introduction to ML and Data Manipulation

Slides based off of Machine Learning at Berkeley
<https://github.com/mlberkeley/Machine-Learning-Decal-Fall-2018>



Overview

Who are we?

What is machine learning?

Class Logistics

General Overview and Context

Machine Learning Pipeline

Python/Numpy/Scikit-Learn

Questions





Who are we?

The Instructors



Frank Ye

U4 Electrical Eng.

- Former Data Science/ML Intern at Splunk (Incoming Full-Time)
- Former Software Engineering Intern at Under Ericsson




Isaac Chan

U4 Electrical Eng.

- Current ML Research Intern at Huawei Technologies
- Former Research Intern at the Graphics & Imaging Lab

The Teaching Assistants (TAs)



Aanika




John




Tiff



Claudia



Jenny



David



Hisham



Meg




Nabil





Ketan



Daoud



Josh



What is Machine Learning?

Age Old Question



Can AI Compose Music



Can AI Paint A Canvas?



Pose Tracking!



Superhuman Reasoning!




Self-Driving Cars




Our View Of Intelligence

Our Distorted View of Intelligence



What Intelligence Is Actually Like





Class Logistics

Goals

Understand major concepts in machine learning

Understand trade-offs between different approaches
(what do I use - when and why?)

- Statistical Learning Methods
- Deep Learning Methods

Gain familiarity to solve ML problems

Develop the skills for your first data
science/machine learning internship

Build a community and have fun!

How we accomplish this

Lectures (8 total)

- 2 hours/week
- Theoretical introduction
- Hands on coding tutorial

Homework (5 total)

- 3-6 hours/week
- Practice implementing material taught in lecture
- Due before the next lecture after being assigned

Office Hours

- 2 hours per week
- Dedicate time for in person support

How we accomplish this (Ct'd)

Final Project

- 3 deliverables + final presentation
- Find a dataset on kaggle and create any real life application using the data (web app, mobile app, IoT, robot, etc)

Final Blog

- Submit on medium at end of course
- Summarize project and reflect on learning

Guest Lectures and Social Events (Optional)

- Gain extra insight from industry professionals
- Work hard and play hard!

Logistical

Join the Slack

<https://slack-link-here.com>

Clone the github

<https://github.com/McGillAISociety/mais-bootcamp-w2019>

Share your Repository

All assignments and deliverables will be marked from your github repository. Learning git is crucial for your career!

Attendance

Attendance is mandatory and will be taken at every lecture.

You may miss up to 1 lecture (there are only 8!) if you have midterms or other commitments.

“Studying for midterms or busy with assignments” is not an excuse. Conflicting schedule with midterm time is excusable.

We have put in a HUGE COMMITMENT to be here for YOU.

If you are not gaining value out of lectures, tell us why!



Evaluation


Homeworks and deliverables are marked based on completion with feedback given to help you improve.

You may have up to 1 incompleteness for homework assignments, but all deliverables must be completed.

Final project and blogs must be completed for certificate.

We aren't paid for teaching and you do not need this for graduation. We all get out what we put into it.

We will not be enforcing plagiarism and we encourage collaboration, because learning together helps AND you get to make new friends!






General Overview And Context




3 different classes of machine learning problems




1- Supervised Learning

Regression



Learning a function for a **continuous** output
Eg. Predicting sales price of house.

Classification



Learning a function for a **categorical** output
Eg. Classifying cats vs dogs in images.



Some Basic Terminology

Features/ Attributes


Target Variable



	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa
7	4.6	3.4	1.4	0.3	setosa
8	5.0	3.4	1.5	0.2	setosa
9	4.4	2.9	1.4	0.2	setosa
10	4.9	3.1	1.5	0.1	setosa
11	5.4	3.7	1.5	0.2	setosa
12	4.8	3.4	1.6	0.2	setosa
13	4.8	3.0	1.4	0.1	setosa
14	4.3	3.0	1.1	0.1	setosa
15	5.8	4.0	1.2	0.2	setosa

2) Unsupervised Learning



3) Reinforcement Learning





Machine Learning Pipeline

The ML Process

1 Identify Problem

Carefully define the problem you want to solve. What specific question are you trying to answer?

2 Gather Data

Figure out what data is needed and where to retrieve it. Does similar data exist or do we need to generate it?

3 Process Data


Format data that can be interpreted by a computer. That includes cleaning, manipulating and extracting important features to feed into the training model.

Train-test Split

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa
7	4.6	3.4	1.4	0.3	setosa
8	5.0	3.4	1.5	0.2	setosa
9	4.4	2.9	1.4	0.2	setosa
10	4.9	3.1	1.5	0.1	setosa
11	5.4	3.7	1.5	0.2	setosa
12	4.8	3.4	1.6	0.2	setosa
13	4.8	3.0	1.4	0.1	setosa
14	4.3	3.0	1.1	0.1	setosa
15	5.8	4.0	1.2	0.2	setosa

Overfitting vs Underfitting

How Overfitting affects Prediction



The ML Process (Continued)

4 Train Model

Training the dataset on your selected model. In practice, datasets are split into train, validation and test sets in order to measure model performance.


5 Evaluate Results

Does the trained model solve your initial problem? Does it satisfy your performance requirements?

6 Repeat!

Improve your model by reiterating the process!

Choosing A Model




K-Fold Cross Validation





Python and Numpy Introduction/Demo





Scikit Learn

Introduction/Demo

Thanks!

Any questions?

Reminders:

Homework 1 and deliverable
1 due before next lecture.

