

McGill **Artificial Intelligence** Society



# Lecture 2: Regression

Slides based off of Machine Learning at Berkeley  
<https://github.com/mlberkeley/Machine-Learning-Decal-Fall-2018>



# Today's Lesson Plan

Linear Regression

Optimization Via Gradient Descent

Logistic Regression

Multinomial Regression

# Recall two main types of supervised algorithms

## Regression:



- Input maps directly to a continuous output space
- Involves estimating or predicting a response
- Ex: predicting housing prices, computing trends in stock market

## Classification:



- Input maps to a class label
- Classification is the act of identifying group membership
- Ex: image classification, semantic classification in text

# Linear Regression

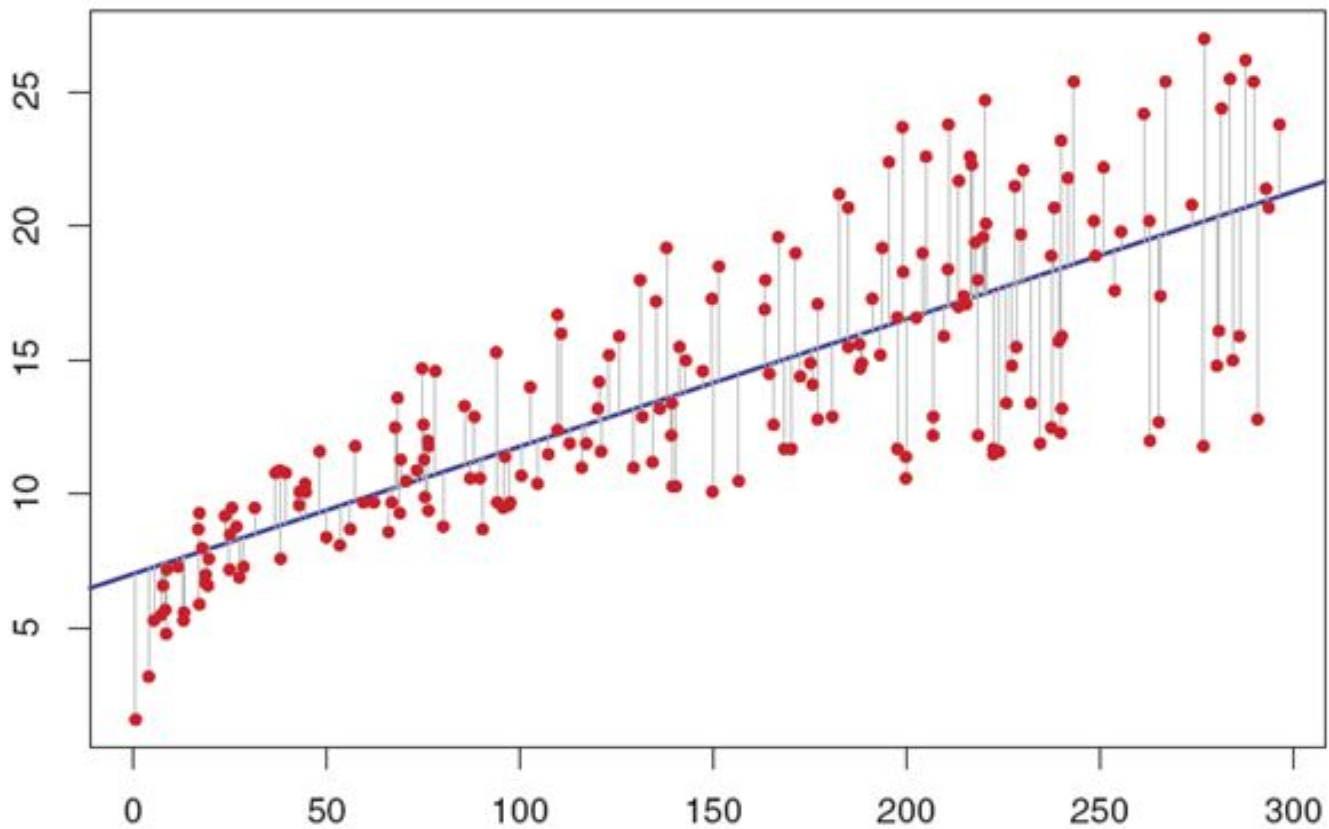


Image credits to Rachna Devasthali from [towarddatascience.com](https://towarddatascience.com)

# Linear Regression

We begin with the general format of a linear regression problem

$$y = \theta_1 x + \theta_0 + \epsilon$$

Where we model real-world practicalities with standard Gaussian noise

$$\epsilon \sim \mathcal{N}(0, \sigma^2)$$

In vector(compact) form, we write:

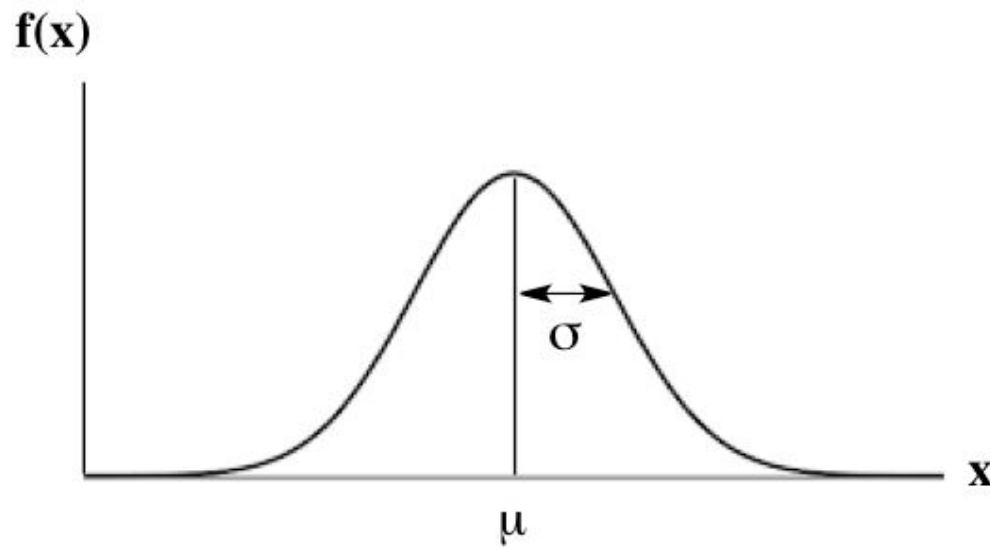
$$y = \theta^T X + \epsilon$$

Probabilistic interpretation:

$$p(y|x, \theta) = \mathcal{N}(y|w^T x, \sigma^2)$$

Quick Review Of Gaussian Distribution!

# The Gaussian Distribution



$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-(x-\mu)^2/2\sigma^2}$$

$$e = 2.71828$$

Distribution defined by its mean, variance



## What are we trying to solve?

In essence, we want to solve the optimization problem

$$\hat{\theta} = \operatorname{argmax}_{\theta} \log p(D|\theta)$$

Where D is a representation of the dataset. The log-likelihood is therefore written as follows:

$$l(\theta) = \log p(D|\theta) = \sum_{i=1}^N \log p(y_i|x_i, \theta)$$

In practice, negative log likelihood is used

$$NLL(\theta) = - \sum_{i=1}^N \log p(y_i|x_i, \theta)$$

# Framing the Optimization Problem

We expand the likelihood equation to its full form

$$l(\theta) = \sum_{i=1}^N \log\left[\left(\frac{1}{2\pi\sigma^2}\right)^{1/2} \exp\left(\frac{-1}{2\sigma^2}(y_i - w^T x_i)^2\right)\right]$$
$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-(x-\mu)^2/2\sigma^2}$$
$$= \frac{-1}{2\sigma^2} RSS(w) - \frac{N}{2} \log(2\pi\sigma^2)$$

Where the residual sum of squares is:

We do not consider this term in our optimization problem as it is constant with respect to the parameters of the model

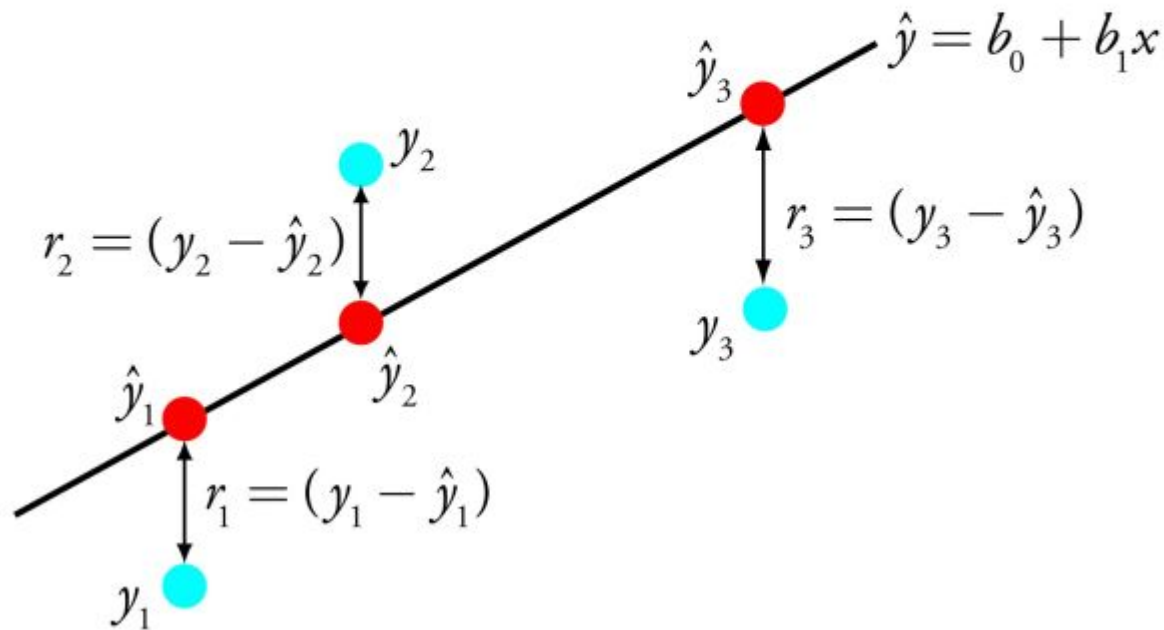
$$RSS(w) = \sum_{i=1}^N (y_i - w^T x_i)^2$$



# Linear Regression -Least Squares Method

$$\text{Let } \hat{y}_i = h(x) = b_0 + b_1 x$$

$$\min J(b_0, b_1)$$



# Linear Regression -Least Squares Method

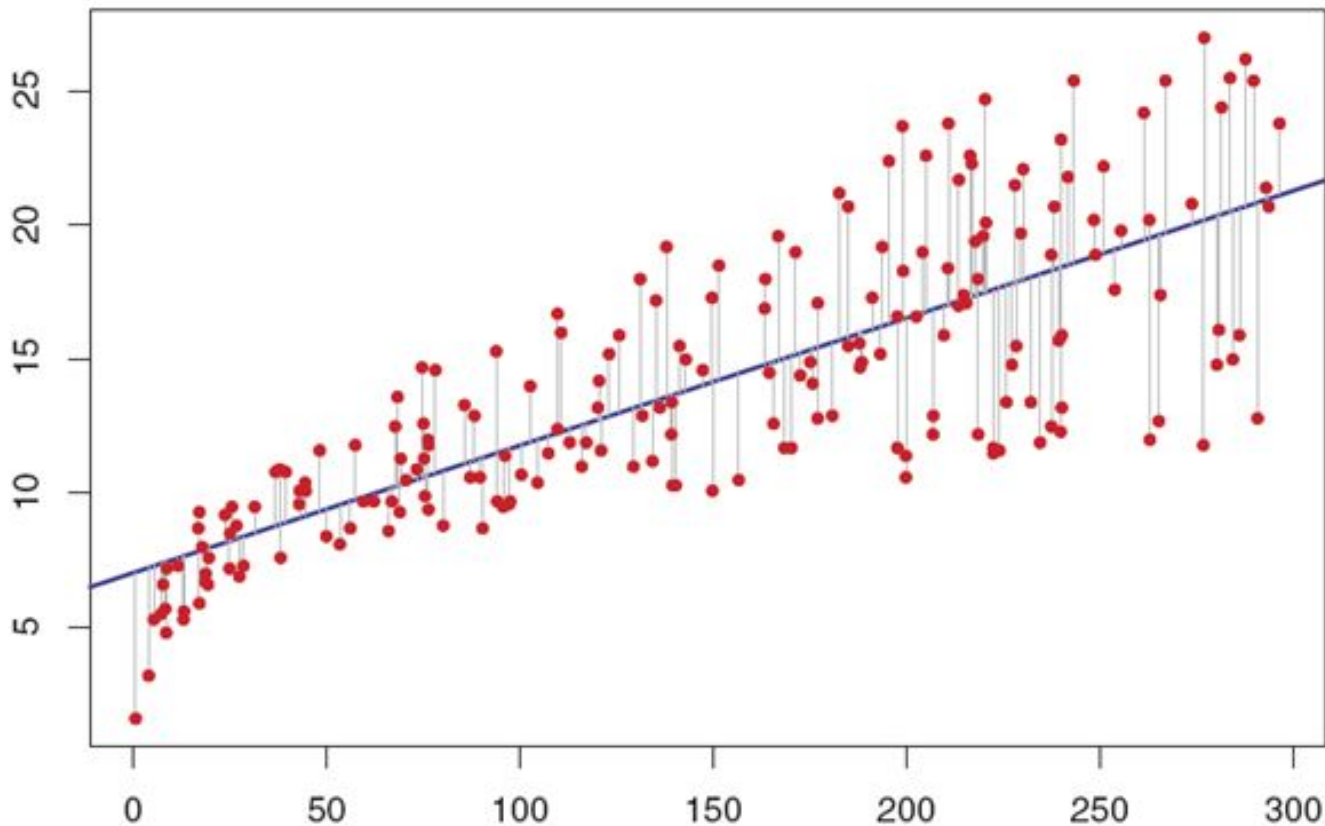
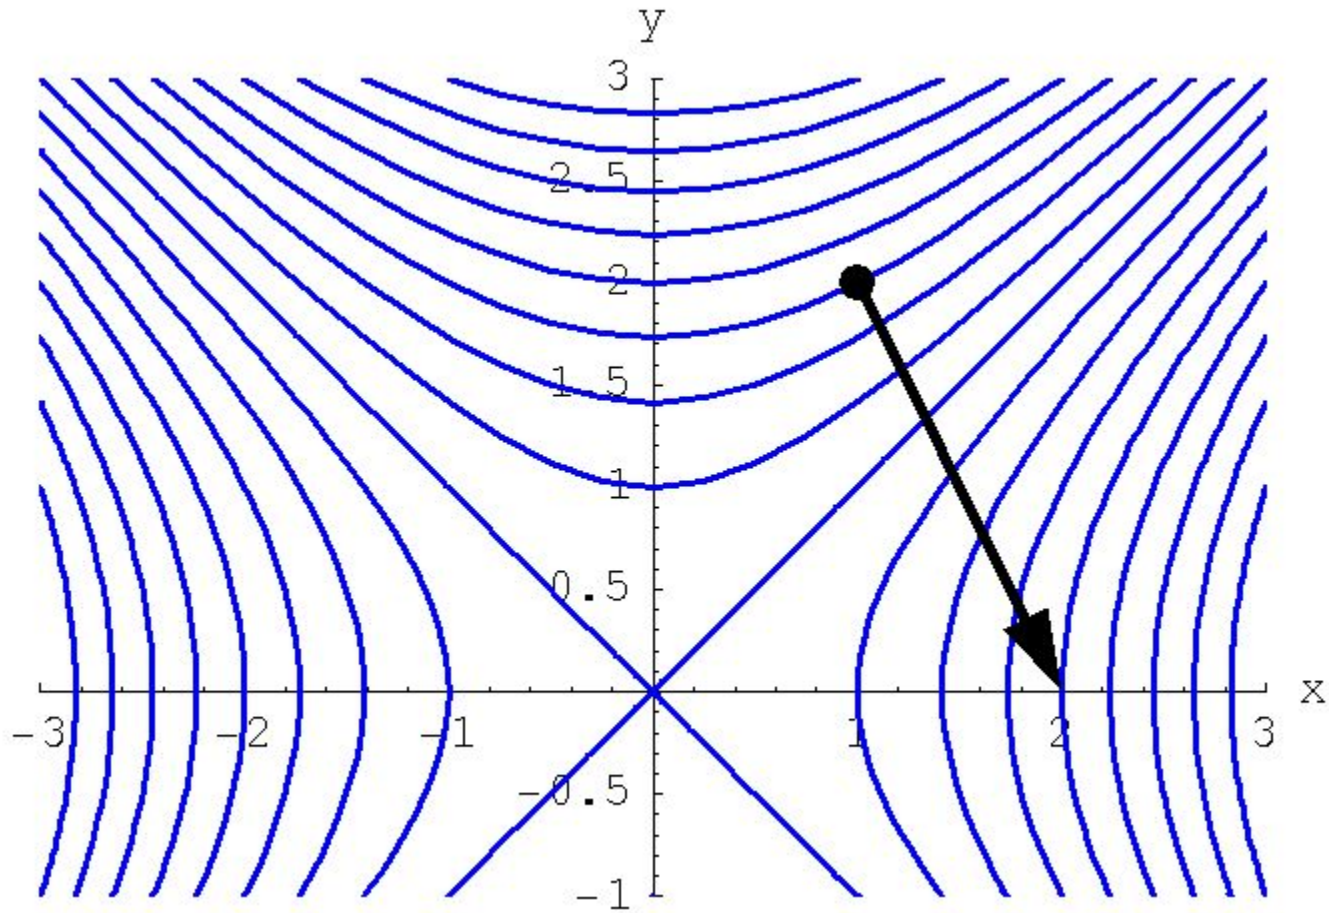


Image credits to Rachna Devasthali from towarddatascience.com

# Optimization Via Gradient Descent



# Derivation of the MLE

Rewriting the objective in a form that is amenable to differentiation

$$NLL(w) = \frac{1}{2}(y - Xw)^T(y - Xw) = \frac{1}{2}w^T(X^T X)w - w^T(X^T y)$$

Where:

$$X^T X = \sum_{i=1}^N x_i x_i^T$$

And XTY follows, we can compute the gradient of the NLL, and we wish to set it to 0:

$$g(w) = [X^T X w - X^T y] = \sum_{i=1}^N x_i (w^T x_i - y_i)$$

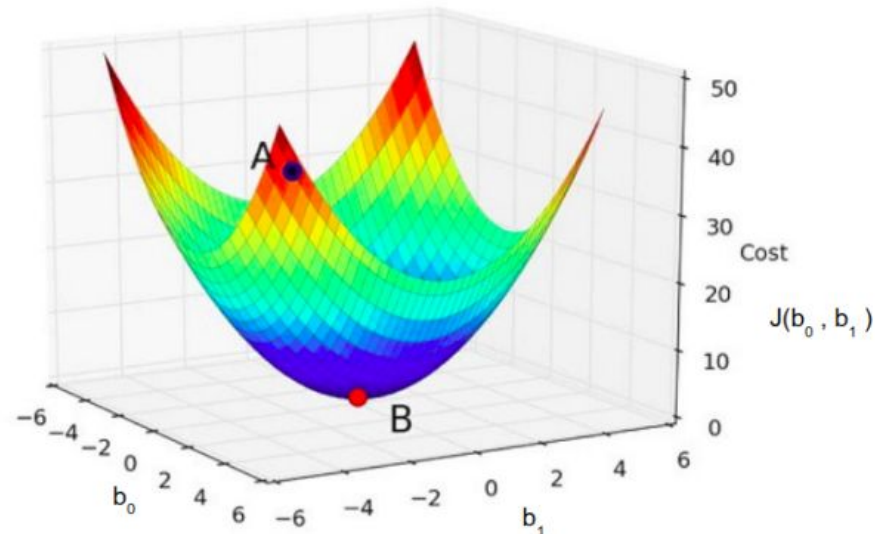
How do we know this reaches a minima?



# Derivation of the MLE

Given that this problem is convex (proof omitted), we know that it has a unique global minimizer, which we can find by setting the gradient to 0, where the solution to the problem is:

$$X^T X w = X^T y$$



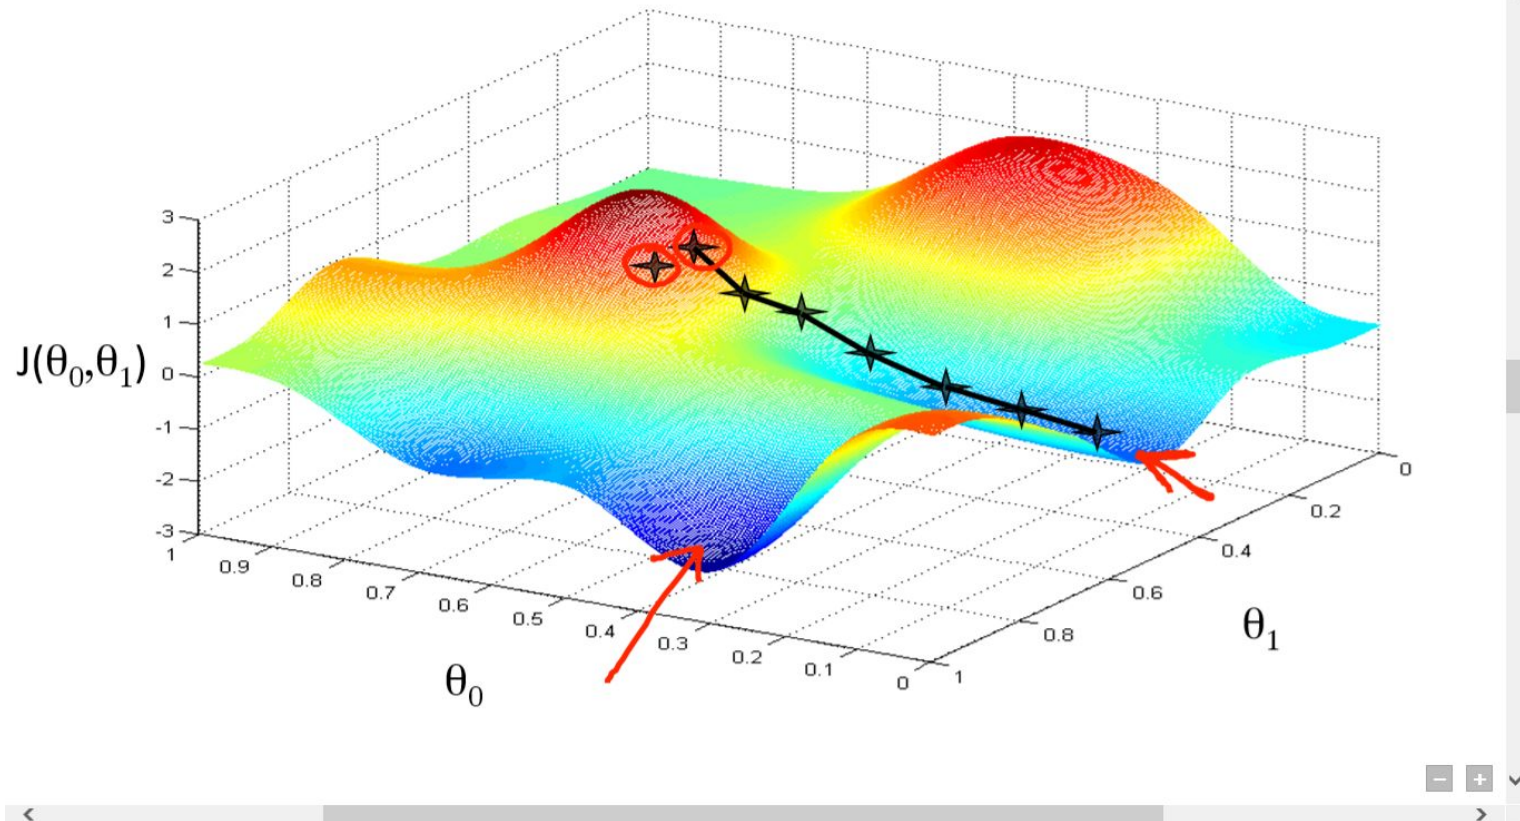
Rearranging the terms, we get the ordinary least squares solution in closed form:

$$\hat{w}_{OLS} = (X^T X)^{-1} X^T y$$

Observe complexity - 3 matrix multiplications, 1 matrix inverse: can compute in polynomial time, **bad for large datasets with many examples, many features.**



# Gradient Descent



# Gradient Descent

We essentially want to iteratively get closer to the minimum of our objective function (defined as the MSE with respect to our weights).

- $w_0, w_1, w_2, \dots$  such that  $MSE(w_0) > MSE(w_1) > MSE(w_2) > \dots$

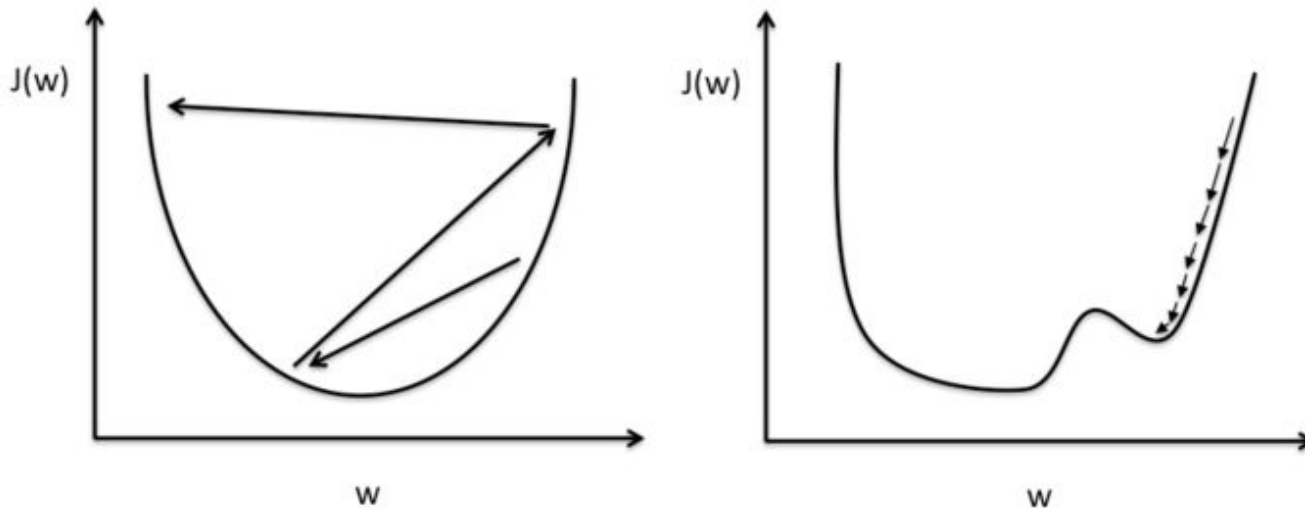
1. Pick initial  $w_0$

2. For  $k = 1, 2, \dots$ ,  
 $w_{k+1} = w_k - \alpha g(w_k)$

where  $\alpha > 0$  is called the “learning rate”

End when  $|w_{k+1} - w_k| < \varepsilon$

## Picking $\alpha$

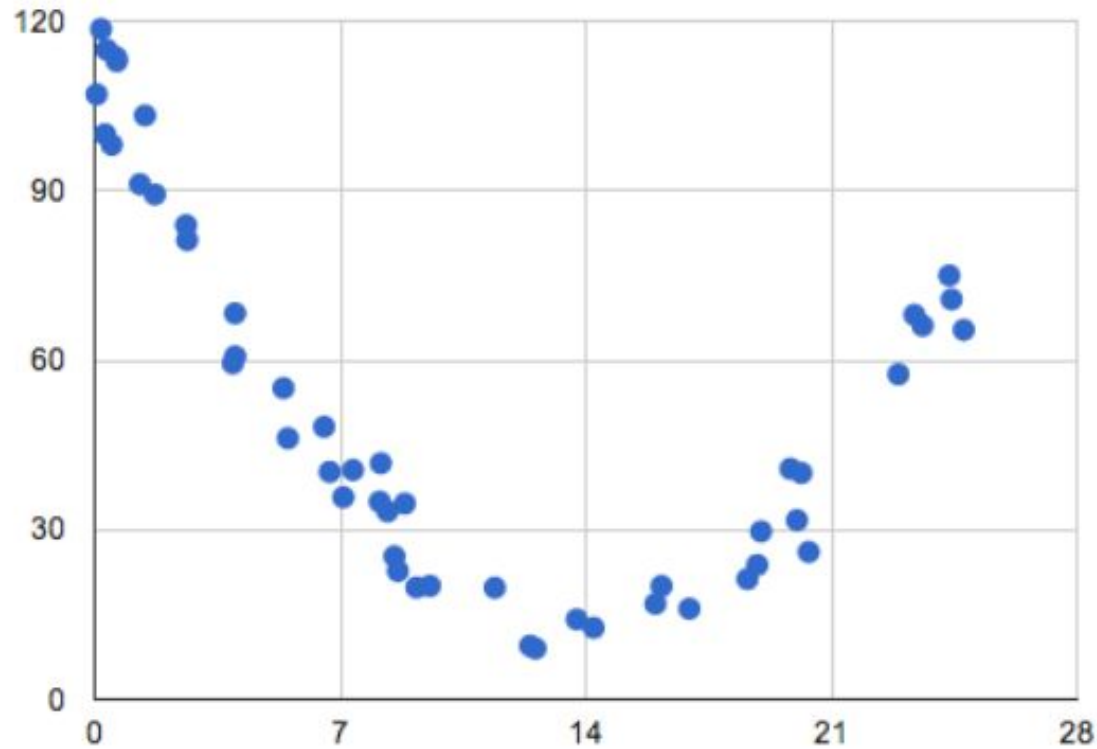


Too large: we “overshoot” and don’t converge to the global minima

Too small: the weights might not move far enough to reach a local minima, slow convergence



## What if Data is Non-Linear?



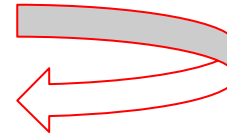
$$h(x) = b_0 + b_1x + b_2x^2$$

# Introducing Polynomial Regression

We linearize our hypothesis  $h(x)$ :

$$h(x) = b_0 + b_1x + b_2x^2$$

$$h(x_1, x_2) = b_0 + b_1x_1 + b_2x_2$$



Such that  $x_1 = x$ , and  $x_2 = x^2$

This way, the model is still **linear** with respect to its **parameters**.

# Introducing Polynomial Regression

Instead of just using  $X$ , we apply a basis function expansion by replacing  $x$  with some non-linear function of the inputs s.t.

$$p(y|x, \theta) = \mathcal{N}(y|w^T \phi(X), \sigma^2)$$

Where

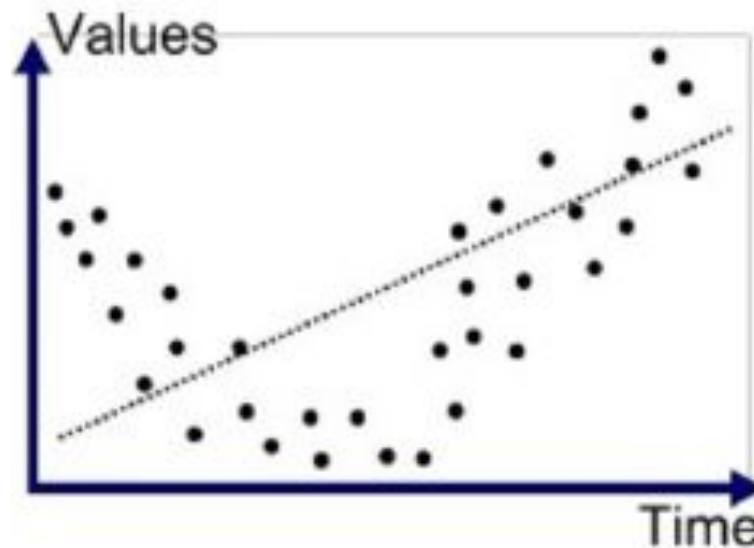
$$\phi(X) = [1, x, x^2, \dots, x^d]$$

:

- model is still linear in parameters  $\mathbf{w}$
- allows to fit nonlinear data ( $\phi(X)$  can be replaced with many other basis function expansions or kernels)

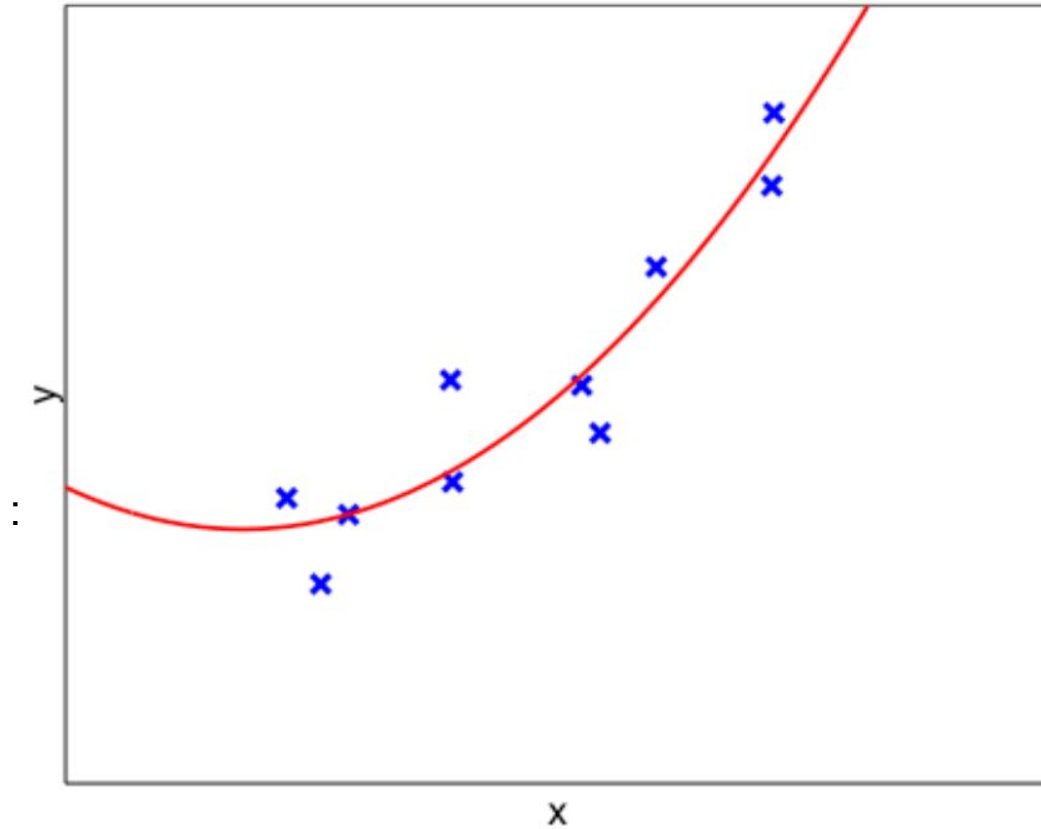
# Overfitting vs. Underfitting

Intuitively, a linear (or a low dimensional polynomial) will not be powerful to fit more complex models → **underfitting**.

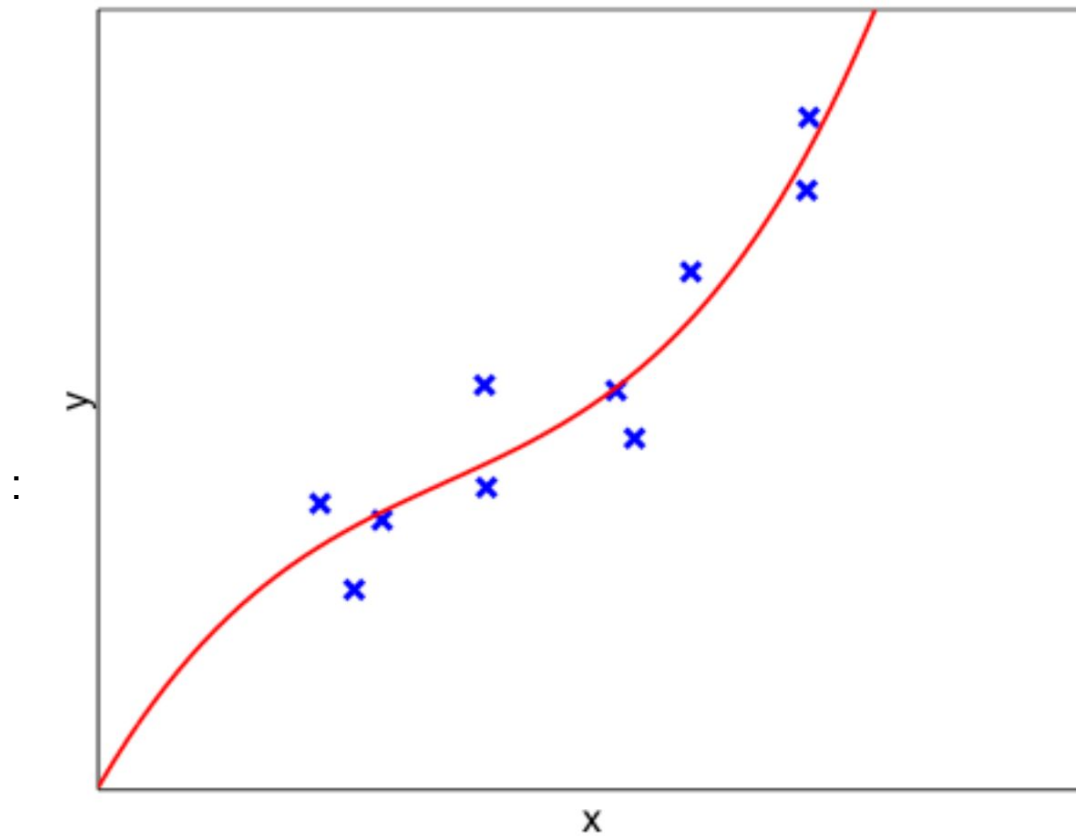


**Overfitting** - the phenomena by which the model is so adapted to the training set that it no longer generalizes well to the underlying model

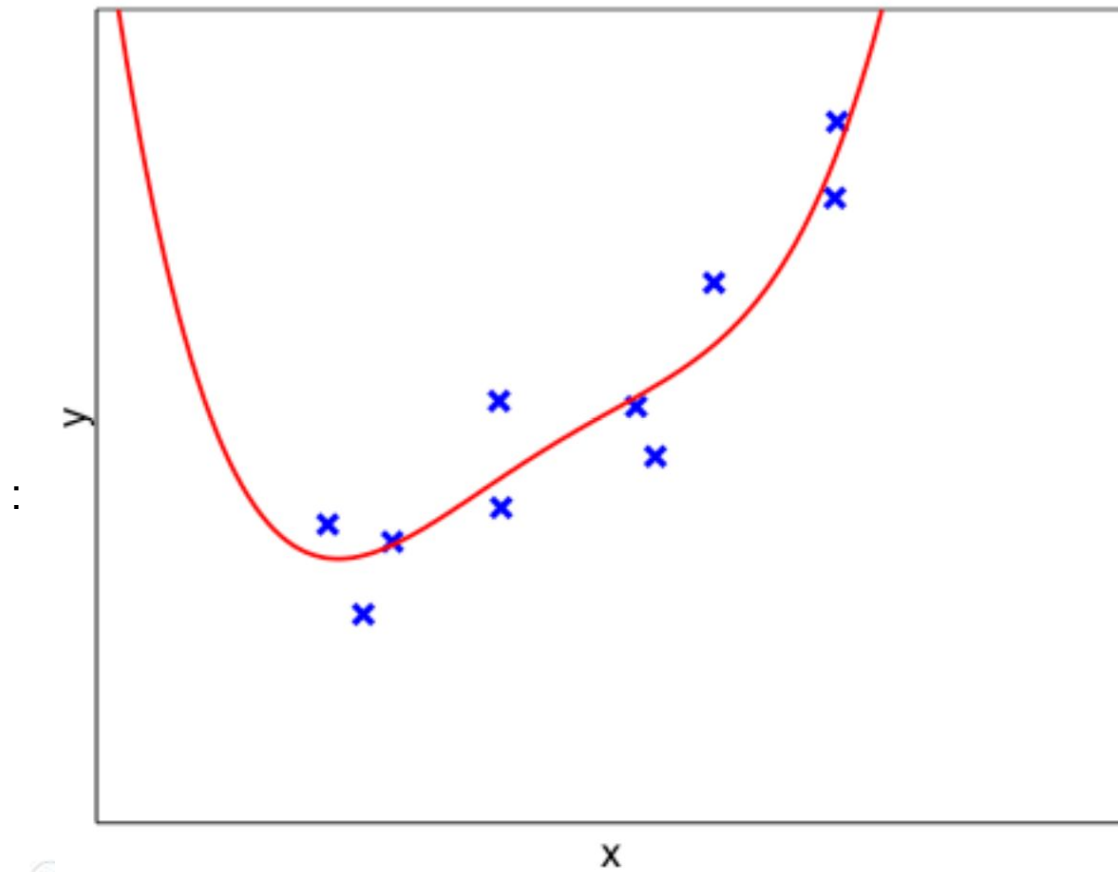
## Order 2 fit



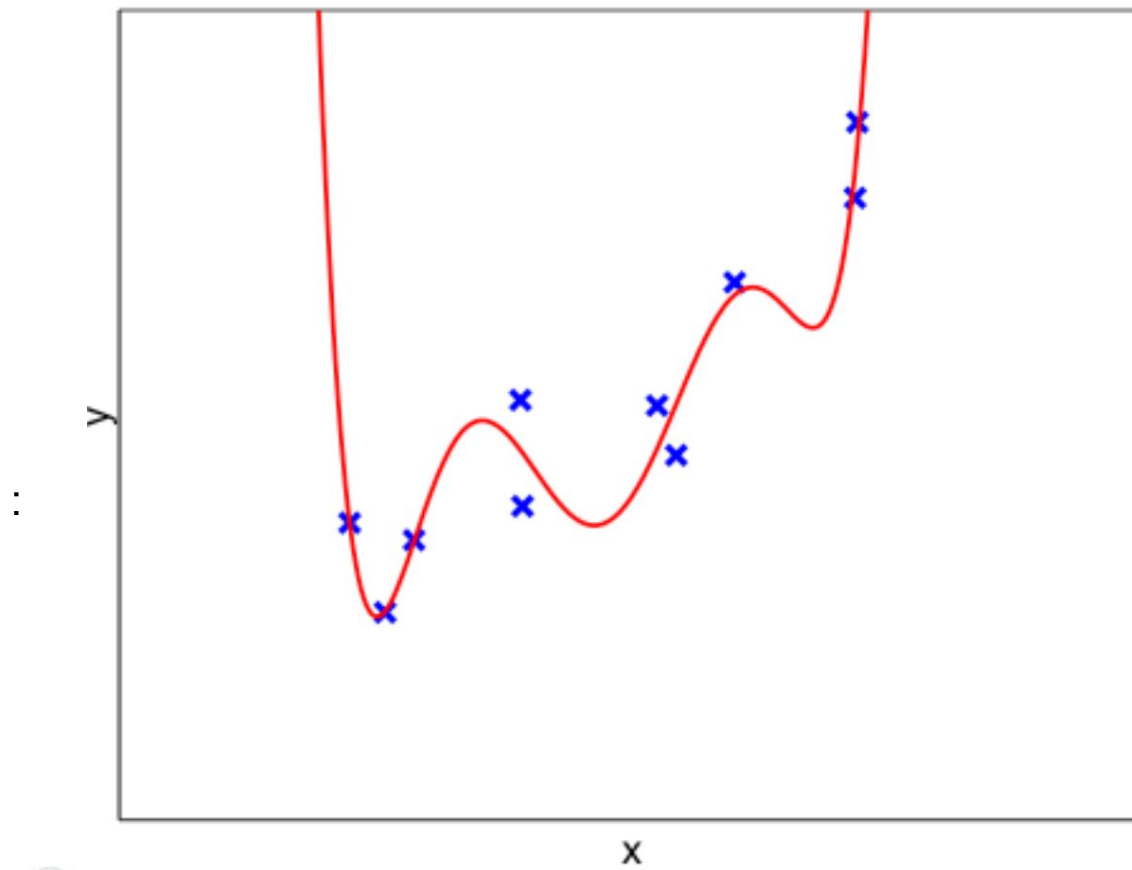
## Order 3 fit



## Order 4 fit

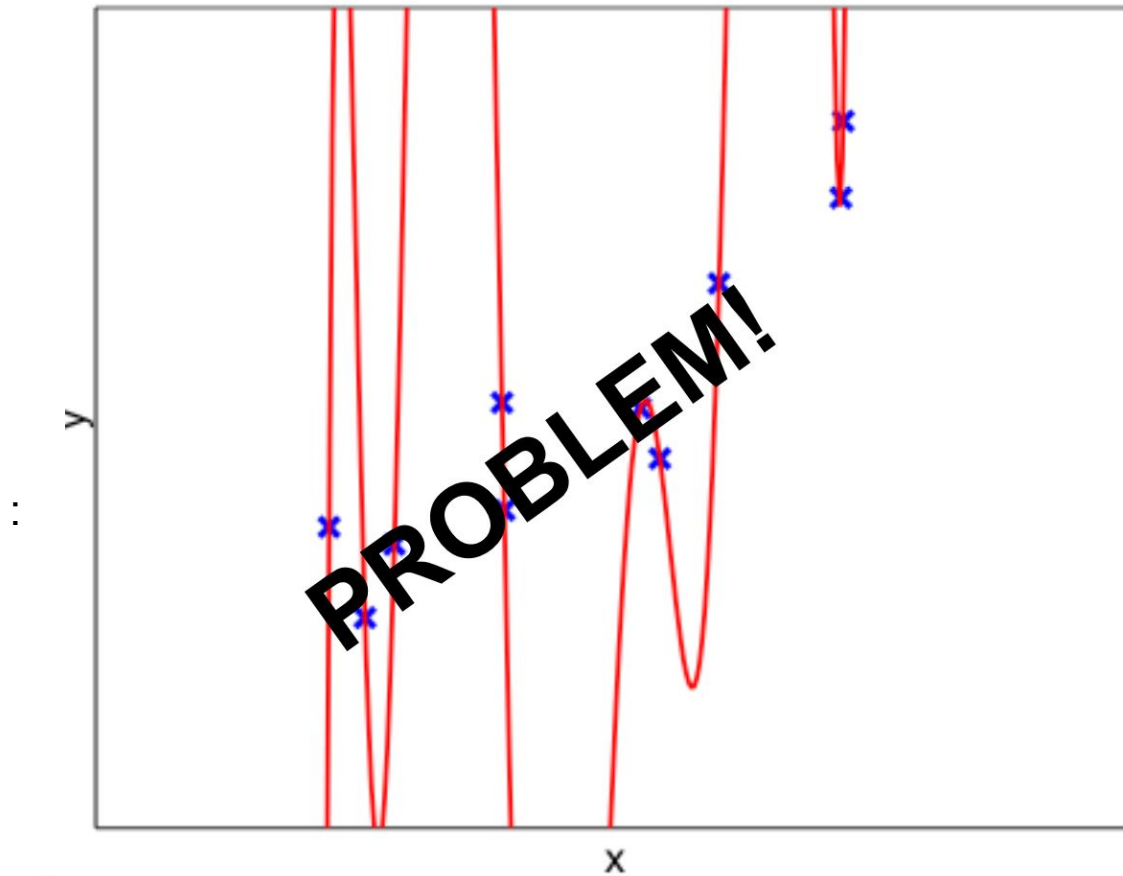


Order 6 fit





Order 9 fit



# Addressing Overfitting


How to address overfitting -

1) **Hyperparameter tuning -**

Simply modify hyperparameters that control the complexity of the model (in this case, the value of  $d$ ) until you get the validation set accuracy optimized

2) **Adding more data-**

- Adding more data, so simply having more training set data can
- allow for a more complex model without overfitting

**More sophisticated methods:** Ridge regression (L2 regularization), dropout in neural networks, lasso regression (L1 regularization), cross-validation, etc. 

:

# 5-MINUTE BREAK

Now Consider the Following Problem...

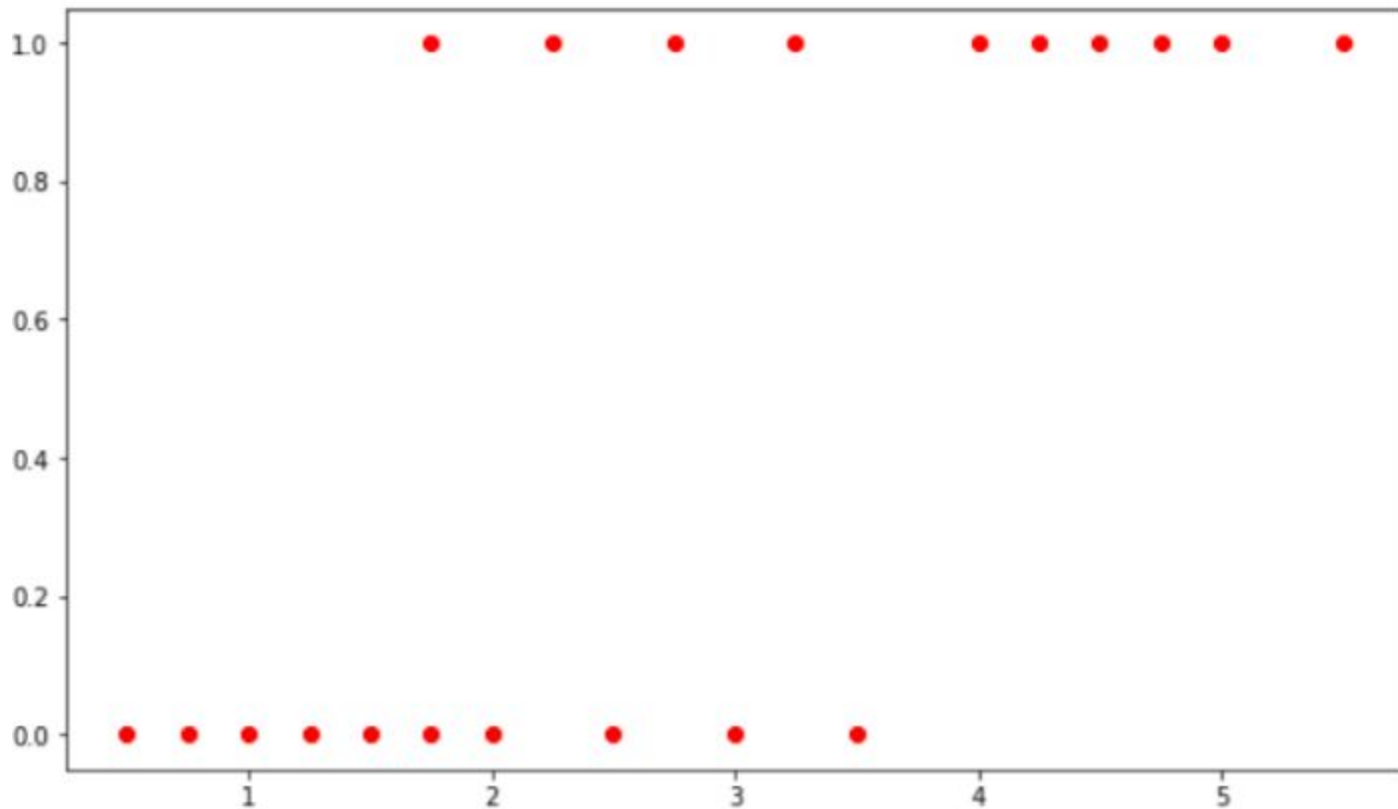
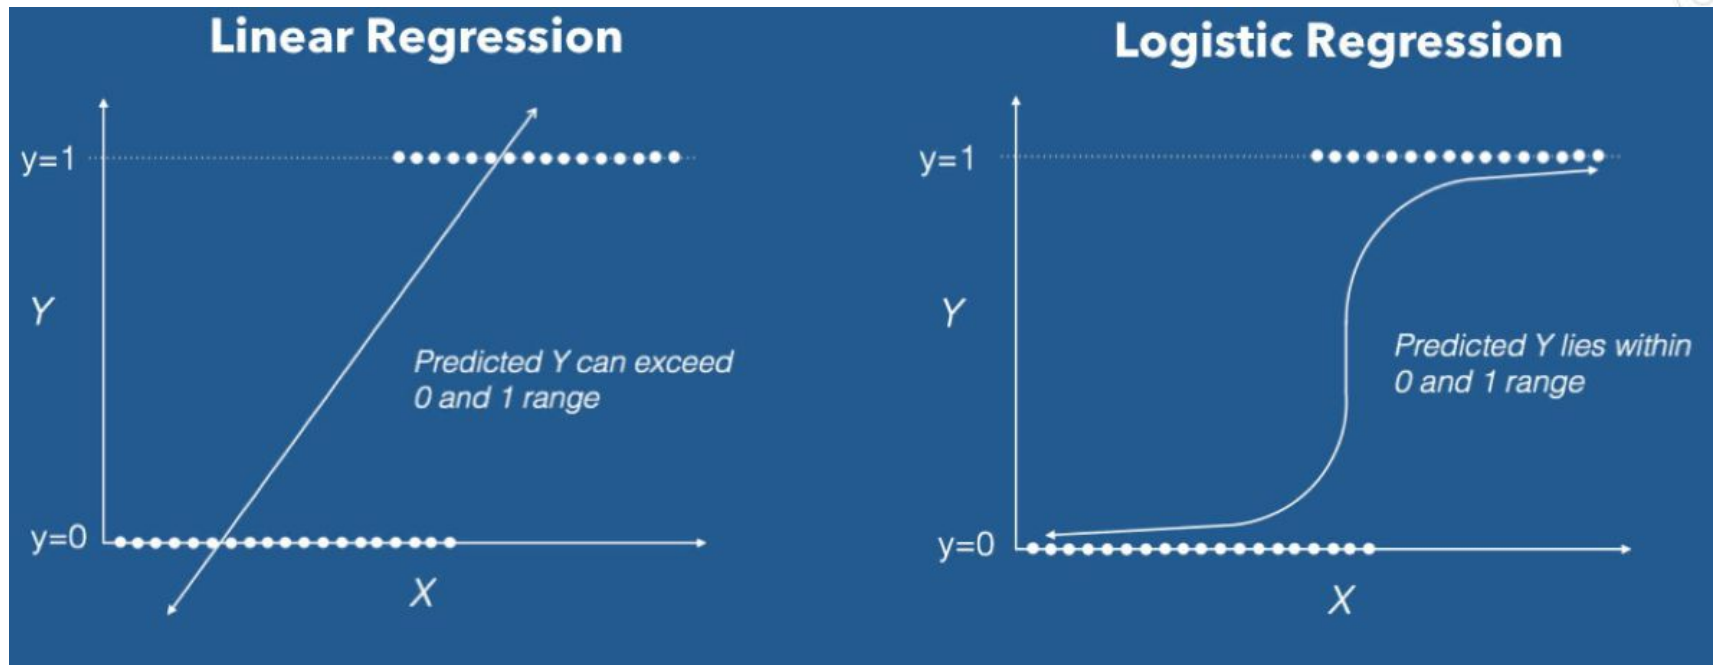


Image credits to Berkeley DeCal Course

# Why Linear Regression Doesn't Work Here



# Logistic Regression

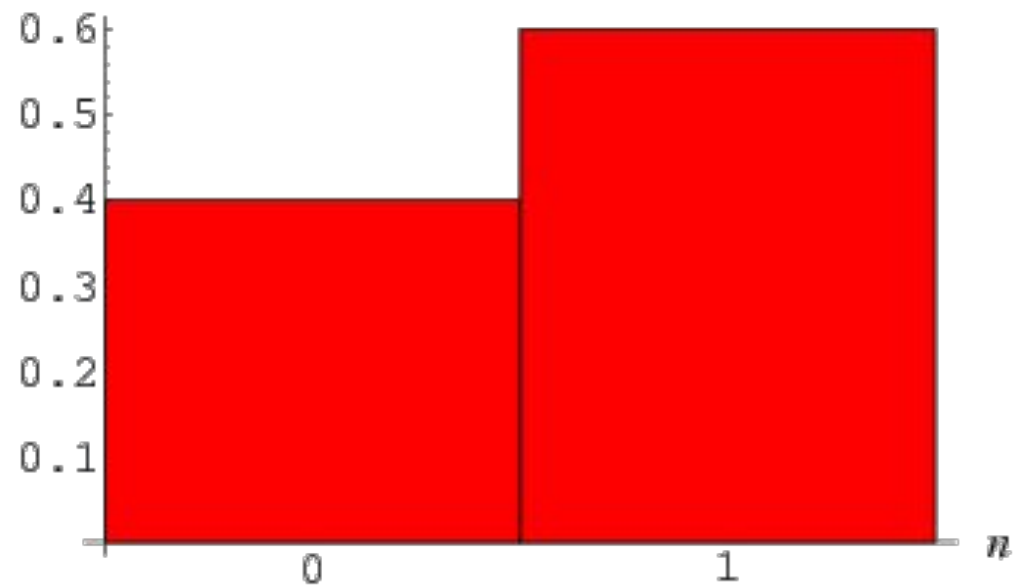
Instead of modeling our response directly, logistic regression models the probability that  $y$  belongs to a certain class:

$$p(y|x, w) = \text{Ber}(y | \text{sigm}(w^T x))$$

Bernoulli, sigmoid... what are those??

# Bernoulli Random Variable

$P(n)$  for  $p = 0.6$



# Sigmoid/ Logistic Function

We want a function  $f$  s.t.  $\text{range}(f) \in [0, 1] \forall X$



## Logistic Function (Ct'd)

$$\begin{aligned} P(X) &= \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}} \\ &= \frac{1}{1 + e^{-(\beta_0 + \beta_1 X)}} \end{aligned}$$

$$\frac{P(X)}{1 - P(X)} = e^{\beta_0 + \beta_1 X}$$

$$\ln \left( \frac{P(X)}{1 - P(X)} \right) = \beta_0 + \beta_1 X$$

The “logit/log-odds” function is linear in X

## Computing Regression Coefficients

Although we can use the least-squares method and estimate using our training data, we prefer the maximum likelihood approach due to its better statistical properties (out of the scope of the course).

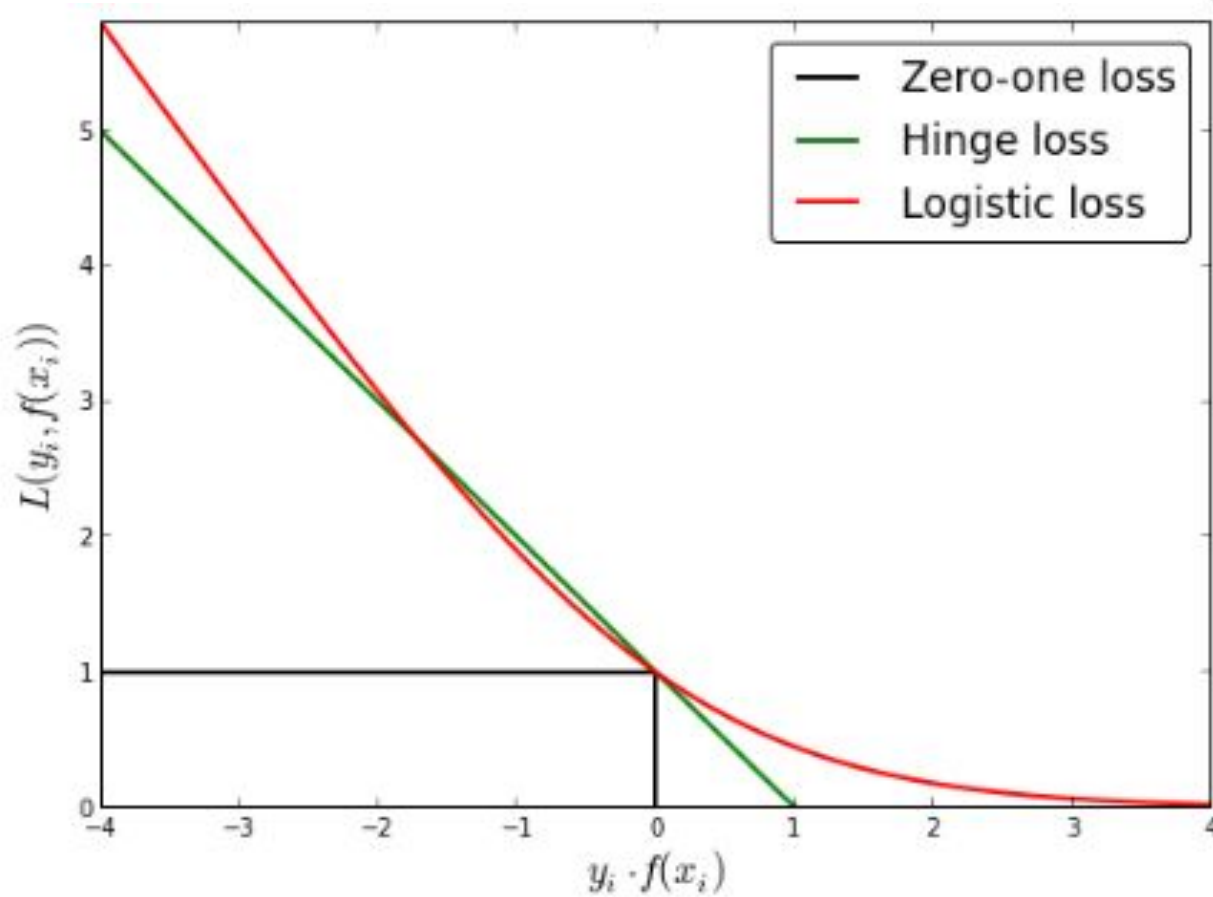
$$\ell(\beta_0, \beta_1) = \prod_{i: y_i=1} p(x_i) \prod_{i': y_{i'}=0} (1 - p(x_{i'})).$$

∴  
We seek to minimize logistic loss:

$$J(b) = - \sum_{i=1}^m \left( y^{(i)} \cdot \ln z^{(i)} + (1 - y^{(i)}) \cdot \ln (1 - z^{(i)}) \right)$$

$$z = h(x) = \frac{1}{1 + e^{-b^T \bar{x}}}$$

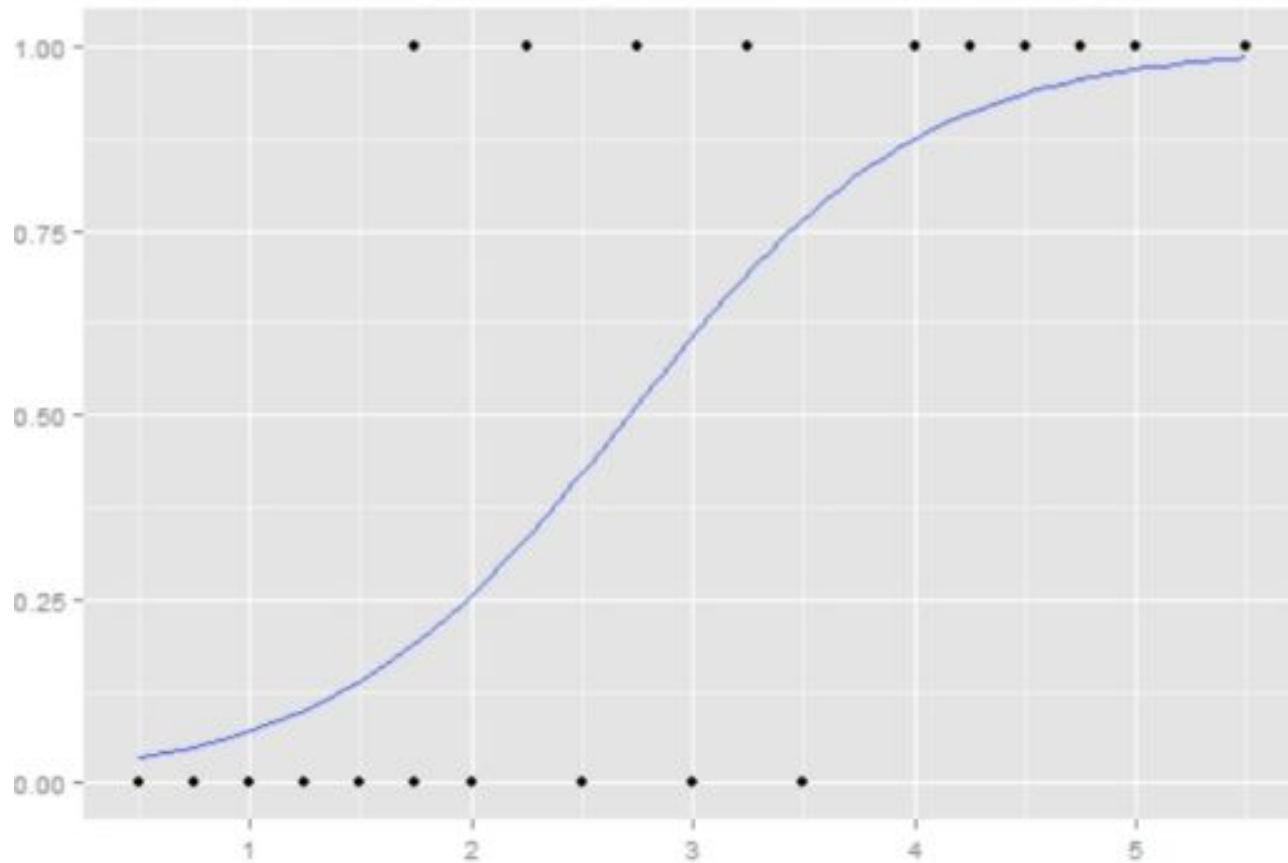
# Logistic Loss



$$J(b) = - \sum_{i=1}^m \left( y^{(i)} \cdot \ln z^{(i)} + (1 - y^{(i)}) \cdot \ln (1 - z^{(i)}) \right)$$

[Image from Stack Exchange](#)

Coefficients should return something like this...



# SUMMARY

	Linear	Logistic
Label Type	Continuous	Categorical
Problem Type	Actual Regression	Actually Classification
Hypothesis	$\theta^T x$	$s(\theta^T x)$
Loss	Mean Squared	Logistic
Analytical Solution	Yes	No

# Thanks!

## Any questions?

Reminders:

Homework 1 and deliverable  
1 due before next lecture.

