

McGill **Artificial Intelligence** Society



# **Lecture 8: Reinforcement Learning: AN Optimization Tour**

A decorative network diagram in the top right corner, featuring a complex web of interconnected nodes and lines, with some nodes highlighted in blue and others in grey.

# Announcements

Last Lecture of the Bootcamp!

Science Fair Coming Up

Blog Posts, make sure they're ready

# Today's Lesson Plan

1. What is Reinforcement Learning
2. Bandit Problems
3. Markov Decision Processes
4. Connection to Supervised Learning
5. Imitation Learning
6. Model-Based Reinforcement Learning
7. Approximate Dynamic Programming
8. Direct Policy Search
  - a. Policy Gradient
  - b. Pure Random Search
  - c. Guided Policy Search

# What is Reinforcement Learning

- So far, we've been concerned with supervised learning
  - Given  $D = \{X, y\}^N$  find function  $f$  such that
$$P(f(X) = Y < \epsilon) > 1 - \delta$$
- Reinforcement Learning concerns itself with a much different paradigm
  - Deals with Autoregressive Time Series Problems
  - Allows us to make long term decisions

# What is Reinforcement Learning


Control Theory



~~Reinforcement Learning~~ is the study of how to use past data to enhance the future manipulation of a dynamical system

Recht 2018

# What is Reinforcement Learning




$\mathbf{s}_t$  – state

$\mathbf{o}_t$  – observation

$\mathbf{a}_t$  – action


$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)$  – policy

$\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)$  – policy (fully observed)



Levine 2018

# Bandit Problems



Barto 2018

# Bandit Problems


- Trying to find  $q_*(a) = \mathbb{E}[R_t | A_t = a]$
- Applying  $\underset{a}{\operatorname{argmax}} Q_t(a) = \mathbb{E}[R_t | A_t = a]$  is called a greedy algorithm
- But what if our estimate of Q is wrong?

Barto 2018



# Exploration vs. Exploitation

- Exploration is the idea of taking what you consider to be suboptimal actions in order to discover if they're optimal
- $\epsilon$  - greedy, with a probability  $\epsilon$  , you take a suboptimal action




Barto 2018

## Further Readings

- Upper Confidence Bounds -  
<http://www.jmlr.org/papers/volume3/auer02a/auer02a.pdf>
- Thompson Sampling Tutorial -  
[https://web.stanford.edu/~bvr/pubs/TS\\_Tutorial.pdf](https://web.stanford.edu/~bvr/pubs/TS_Tutorial.pdf)
- E3 algorithm -  
<https://www.cis.upenn.edu/~mkearns/papers/reinforcement.pdf>
- The “Bandit” book: <https://banditalgs.com/>

# Markov Decision Processes

- Given,  $s_t$  take action  $a_t$ , and get tuple  $(r_t, s_{t+1})$
- Goal is to maximize  $V(s_0) = \sum_{i=1}^N r_t$



Barto 2018

# Markov Decision Processes

- Problem with this formulation
- Say you're solving a maze, and by solving the maze, you get a reward = 1, otherwise, you get reward = 0
  - Previous formulation would therefore not care if you completed it in the next step or in 10 years
- More relatably
  - 1 million dollars is worth more to you tomorrow than it is in 10 years
  - But how do we capture this?

Barto 2018

# Markov Decision Processes

- Instead of using  $V(s_0) = \sum_{i=1}^N r_t$ , we solve the discounted sum of rewards

$$V(s_0) = \sum_{i=1}^N \gamma^t r_t$$

- A morbid interpretation of this is that there is a probability gamma at each time step that you will die, and therefore, you could interpret the formulation as
  - With probability gamma, you die
  - With probability 1 - gamma, you get reward  $r_t$
- Full formulation

Barto 2018

$$\begin{aligned} &\text{maximize} && \mathbb{E}_{e_t}[\sum_{t=0}^N R_t[x_t, u_t]] \\ &\text{subject to} && x_{t+1} = f_t(x_t, u_t, e_t) \\ &&& (x_0 \text{ given}). \end{aligned}$$

# Connection to Supervised Learning

- Instead of  $x \rightarrow y$ , we now have to consider two variables: action  $a$  (or  $u$ ) and reward  $r$
- Reinforcement Learning is evidently more challenging than supervised learning, but it allows us to work in a much more complex paradigm.
- Difficulties:
  - Non i.i.d data
  - Complex feedback loop
  - Exploration Exploitation Tradeoff
  - Online updates
  - Defining rewards

Recht 2018

# Imitation Learning



$\mathbf{o}_t$   
 $\mathbf{a}_t$




supervised  
learning

$\pi_{\theta}(\mathbf{a}_t|\mathbf{o}_t)$

behavior cloning

Levine 2018


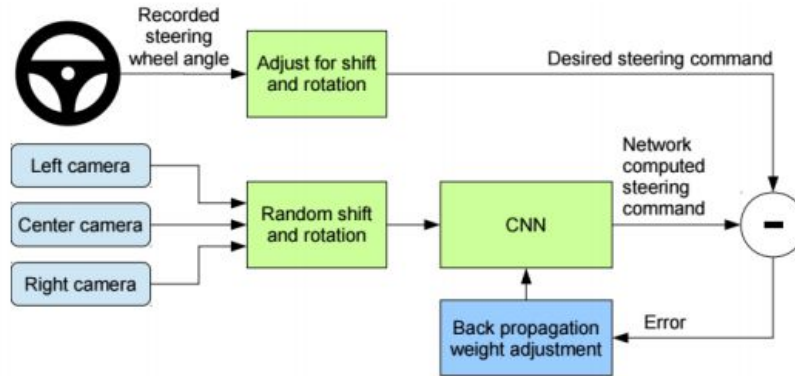
# Does It Work? Not really



Levine 2018



# Does It Work? Ok...maybe sometimes



Levine 2018


# A Slightly Smarter Solution

## DAgger: Dataset Aggregation

goal: collect training data from  $p_{\pi_\theta}(\mathbf{o}_t)$  instead of  $p_{\text{data}}(\mathbf{o}_t)$


how? just run  $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$

but need labels  $\mathbf{a}_t$ !

- 
1. train  $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$  from human data  $\mathcal{D} = \{\mathbf{o}_1, \mathbf{a}_1, \dots, \mathbf{o}_N, \mathbf{a}_N\}$
  2. run  $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$  to get dataset  $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
  3. Ask human to label  $\mathcal{D}_\pi$  with actions  $\mathbf{a}_t$
  4. Aggregate:  $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

Levine 2018

# Why We Might Fail To Fit

- 
1. Non-Markovian behavior
  2. Multimodal behavior

$$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)$$

behavior depends only  
on current observation

$$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_1, \dots, \mathbf{o}_t)$$

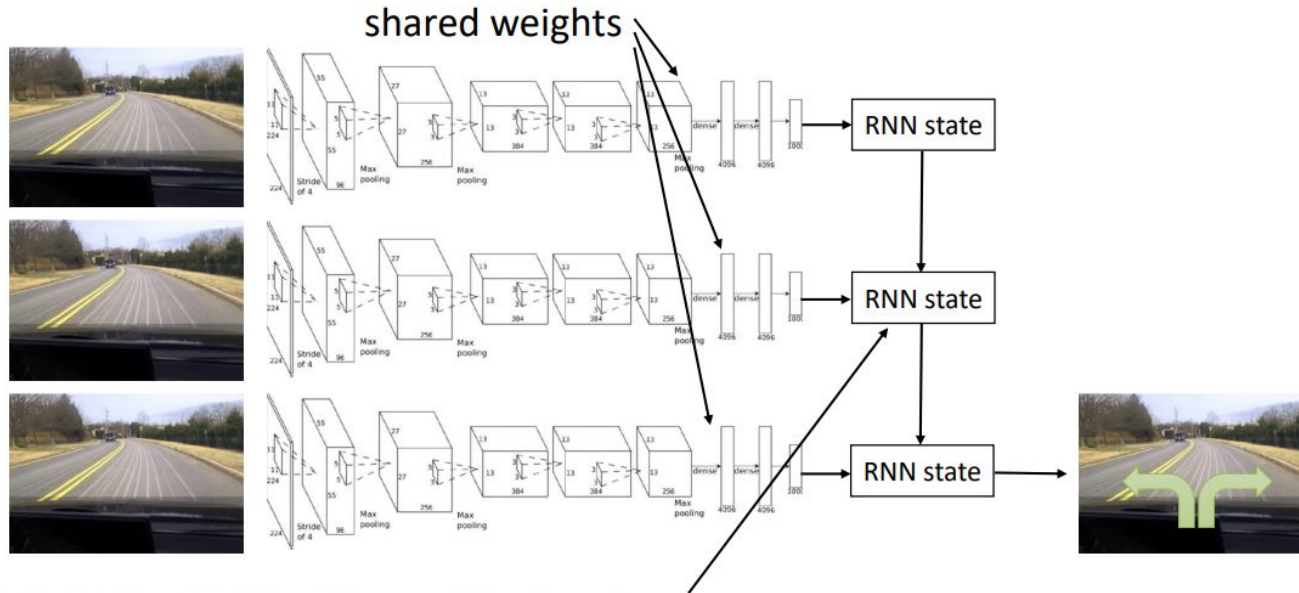
behavior depends on  
all past observations

If we see the same thing  
twice, we do the same thing  
twice, regardless of what  
happened before

Often very unnatural for  
human demonstrators

Levine 2018

# How We Could Make Use Of Full History



Typically, LSTM cells work better here



Levine 2018

# How We Could Make Use Of Full History

Why might we fail to fit the expert?

1. Non-Markovian behavior

→ 2. Multimodal behavior




1. Output mixture of Gaussians
2. Latent variable models
3. Autoregressive discretization



Levine 2018

# Imitation Learning: Further Reading

- Generative Adversarial Imitation Learning (GAIL):  
<https://arxiv.org/abs/1606.03476>
- Learning Robust Rewards with Adversarial Inverse Reinforcement Learning: <https://arxiv.org/abs/1710.11248>



# Model-Based RL

- Fit a Predictive Model and use Dynamic Programming to plan for prescribed control problem
  - This estimated model is known as the nominal model
  - Control design is nominal controller
- Estimation of a dynamical system is called system identification
- Suppose we want to build a predictor of  $x_{t+1}$  from past history data
  - Simple strategy is to inject random probing sequence  $\{u\}^t$  and see how the system responds. Up to stochastic noise, we should get

$$x_{t+1} \approx \varphi(x_t, u_t),$$

Recht 2018

# Model-Based RL

- In such a case, we are focused on the following objective function

$$\text{minimize}_{\varphi} \quad \sum_{t=0}^{N-1} \|x_{t+1} - \varphi(x_t, u_t)\|^2.$$

- So the optimal control problem becomes

$$\begin{aligned} &\text{maximize} \quad \mathbb{E}_{\omega_t} [\sum_{t=0}^N R(x_t, u_t)] \\ &\text{subject to} \quad x_{t+1} = \hat{\varphi}(x_t, u_t) + \omega_t, \quad u_t = \pi_t(\tau_t). \end{aligned}$$



# Model-Based RL

---

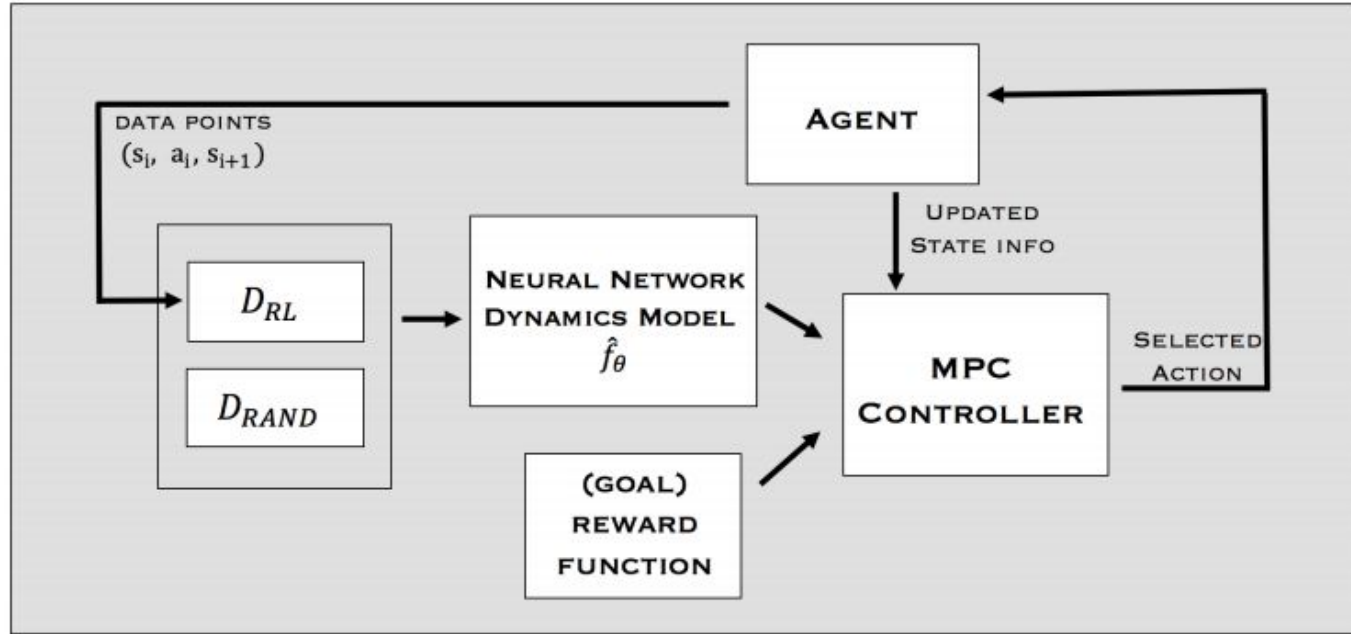
**Algorithm 1** Model-based Reinforcement Learning

---

- 1: gather dataset  $\mathcal{D}_{\text{RAND}}$  of random trajectories
  - 2: initialize empty dataset  $\mathcal{D}_{\text{RL}}$ , and randomly initialize  $\hat{f}_{\theta}$
  - 3: **for** iter=1 **to** max\_iter **do**
  - 4:   train  $\hat{f}_{\theta}(\mathbf{s}, \mathbf{a})$  by performing gradient descent on Eqn. 2,  
    using  $\mathcal{D}_{\text{RAND}}$  and  $\mathcal{D}_{\text{RL}}$
  - 5:   **for**  $t = 1$  **to**  $T$  **do**
  - 6:     get agent's current state  $\mathbf{s}_t$
  - 7:     use  $\hat{f}_{\theta}$  to estimate optimal action sequence  $\mathbf{A}_t^{(H)}$   
      (Eqn. 4)
  - 8:     execute first action  $\mathbf{a}_t$  from selected action sequence  
       $\mathbf{A}_t^{(H)}$
  - 9:     add  $(\mathbf{s}_t, \mathbf{a}_t)$  to  $\mathcal{D}_{\text{RL}}$
  - 10:   **end for**
  - 11: **end for**
- 

Nagabandi 2017

# Model-Based RL



Nagabandi 2017

## Case Study: Linear Quadratic Regulator

$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = \mathcal{N}(f(\mathbf{x}_t, \mathbf{u}_t), \Sigma)$$

$$f(\mathbf{x}_t, \mathbf{u}_t) \approx \mathbf{A}_t \mathbf{x}_t + \mathbf{B}_t \mathbf{u}_t$$

$$\mathbf{A}_t = \frac{df}{d\mathbf{x}_t} \quad \mathbf{B}_t = \frac{df}{d\mathbf{u}_t}$$

$Q(\mathbf{x}_t, \mathbf{u}_t)$  is the cost to go: total cost we get after taking an action

$$Q(\mathbf{x}_t, \mathbf{u}_t) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{Q}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{q}_t$$

## Case Study: Linear Quadratic Regulator

$$\begin{aligned} \text{minimize} \quad & \mathbb{E}_{e_t} \left[ \frac{1}{2} \sum_{t=0}^N x_t^T Q x_t + u_t^T R u_t + \frac{1}{2} x_{N+1}^T S x_{N+1} \right], \\ \text{subject to} \quad & x_{t+1} = A x_t + B u_t + e_t, \quad u_t = \pi_t(\tau_t) \\ & (x_0 \text{ given}). \end{aligned}$$

$$u_t = -K x_t \quad K = (R + B^T M B)^{-1} B^T M A$$

$$M = Q + A^T M A - (A^T M B)(R + B^T M B)^{-1} (B^T M A)$$

But we don't actually know the model!

# Case Study: Linear Quadratic Regulator

$$\text{minimize}_{A,B} \sum_{t=0}^{N-1} \|x_{t+1} - Ax_t - Bu_t\|^2.$$

1. Use supervised learning to learn a coarse model of the dynamical system to be controlled. I will refer to the system estimate as the *nominal system*.
2. Using either prior knowledge or statistical tools like the bootstrap, build probabilistic guarantees about the distance between the nominal system and the true, unknown dynamics.
3. Solve a *robust optimization* problem that optimizes control of the nominal system while penalizing signals with respect to the estimated uncertainty, ensuring stable, robust execution.


$$\frac{\hat{J} - J_{\star}}{J_{\star}} = \tilde{O} \left( \sqrt{\frac{d+p}{T}} \right).$$

State dimension  $d$ , control dimension  $p$ ,  
trajectory  $T$

# Guided Policy Search

shooting method: optimize over actions only


$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} c(\mathbf{x}_1, \mathbf{u}_1) + c(f(\mathbf{x}_1, \mathbf{u}_1), \mathbf{u}_2) + \dots + c(f(f(\dots)), \mathbf{u}_T)$$



# Known Model Methods

collocation method: optimize over actions and states, with constraints

$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T, \mathbf{x}_1, \dots, \mathbf{x}_T} \sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t) \text{ s.t. } \mathbf{x}_t = f(\mathbf{x}_{t-1}, \mathbf{u}_{t-1})$$



# KMM: Start with Linear Dynamics

$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} c(\mathbf{x}_1, \mathbf{u}_1) + c(f(\mathbf{x}_1, \mathbf{u}_1), \mathbf{u}_2) + \dots + c(f(f(\dots)), \mathbf{u}_T)$$

$$f(\mathbf{x}_t, \mathbf{u}_t) = \mathbf{F}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \mathbf{f}_t$$

linear

$$c(\mathbf{x}_t, \mathbf{u}_t) = \frac{1}{2} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{C}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{c}_t$$

quadratic



# KMM: Start with Linear Dynamics

Base case: solve for  $\mathbf{u}_T$  *only*

$$Q(\mathbf{x}_T, \mathbf{u}_T) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_T \\ \mathbf{u}_T \end{bmatrix}^T \mathbf{C}_T \begin{bmatrix} \mathbf{x}_T \\ \mathbf{u}_T \end{bmatrix} + \begin{bmatrix} \mathbf{x}_T \\ \mathbf{u}_T \end{bmatrix}^T \mathbf{c}_T$$

$$\nabla_{\mathbf{u}_T} Q(\mathbf{x}_T, \mathbf{u}_T) = \mathbf{C}_{\mathbf{u}_T, \mathbf{x}_T} \mathbf{x}_T + \mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T} \mathbf{u}_T + \mathbf{c}_{\mathbf{u}_T}^T = 0$$

$$\mathbf{u}_T = -\mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T}^{-1} (\mathbf{C}_{\mathbf{u}_T, \mathbf{x}_T} \mathbf{x}_T + \mathbf{c}_{\mathbf{u}_T})$$

$$\mathbf{u}_T = \mathbf{K}_T \mathbf{x}_T + \mathbf{k}_T$$

$$\mathbf{K}_T = -\mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T}^{-1} \mathbf{C}_{\mathbf{u}_T, \mathbf{x}_T}$$

$$\mathbf{k}_T = -\mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T}^{-1} \mathbf{c}_{\mathbf{u}_T}$$

$$\mathbf{C}_T = \begin{bmatrix} \mathbf{C}_{\mathbf{x}_T, \mathbf{x}_T} & \mathbf{C}_{\mathbf{x}_T, \mathbf{u}_T} \\ \mathbf{C}_{\mathbf{u}_T, \mathbf{x}_T} & \mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T} \end{bmatrix}$$

$$\mathbf{c}_T = \begin{bmatrix} \mathbf{c}_{\mathbf{x}_T} \\ \mathbf{c}_{\mathbf{u}_T} \end{bmatrix}$$

# KMM: Start with Linear Dynamics

$$\mathbf{u}_T = \mathbf{K}_T \mathbf{x}_T + \mathbf{k}_T$$

$$\mathbf{K}_T = -\mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T}^{-1} \mathbf{C}_{\mathbf{u}_T, \mathbf{x}_T}$$

$$\mathbf{k}_T = -\mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T}^{-1} \mathbf{c}_{\mathbf{u}_T}$$

$$Q(\mathbf{x}_T, \mathbf{u}_T) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_T \\ \mathbf{u}_T \end{bmatrix}^T \mathbf{C}_T \begin{bmatrix} \mathbf{x}_T \\ \mathbf{u}_T \end{bmatrix} + \begin{bmatrix} \mathbf{x}_T \\ \mathbf{u}_T \end{bmatrix}^T \mathbf{c}_T$$

Since  $\mathbf{u}_T$  is fully determined by  $\mathbf{x}_T$ , we can eliminate it via substitution!

$$V(\mathbf{x}_T) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_T \\ \mathbf{K}_T \mathbf{x}_T + \mathbf{k}_T \end{bmatrix}^T \mathbf{C}_T \begin{bmatrix} \mathbf{x}_T \\ \mathbf{K}_T \mathbf{x}_T + \mathbf{k}_T \end{bmatrix} + \begin{bmatrix} \mathbf{x}_T \\ \mathbf{K}_T \mathbf{x}_T + \mathbf{k}_T \end{bmatrix}^T \mathbf{c}_T$$

$$V(\mathbf{x}_T) = \frac{1}{2} \mathbf{x}_T^T \mathbf{C}_{\mathbf{x}_T, \mathbf{x}_T} \mathbf{x}_T + \frac{1}{2} \mathbf{x}_T^T \mathbf{C}_{\mathbf{x}_T, \mathbf{u}_T} \mathbf{K}_T \mathbf{x}_T + \frac{1}{2} \mathbf{x}_T^T \mathbf{K}_T^T \mathbf{C}_{\mathbf{u}_T, \mathbf{x}_T} \mathbf{x}_T + \frac{1}{2} \mathbf{x}_T^T \mathbf{K}_T^T \mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T} \mathbf{K}_T \mathbf{x}_T + \mathbf{x}_T^T \mathbf{K}_T^T \mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T} \mathbf{k}_T + \frac{1}{2} \mathbf{x}_T^T \mathbf{C}_{\mathbf{x}_T, \mathbf{u}_T} \mathbf{k}_T + \mathbf{x}_T^T \mathbf{c}_{\mathbf{x}_T} + \mathbf{x}_T^T \mathbf{K}_T^T \mathbf{c}_{\mathbf{u}_T} + \text{const}$$

$$V(\mathbf{x}_T) = \text{const} + \frac{1}{2} \mathbf{x}_T^T \mathbf{V}_T \mathbf{x}_T + \mathbf{x}_T^T \mathbf{v}_T$$

$$\mathbf{V}_T = \mathbf{C}_{\mathbf{x}_T, \mathbf{x}_T} + \mathbf{C}_{\mathbf{x}_T, \mathbf{u}_T} \mathbf{K}_T + \mathbf{K}_T^T \mathbf{C}_{\mathbf{u}_T, \mathbf{x}_T} + \mathbf{K}_T^T \mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T} \mathbf{K}_T$$

$$\mathbf{v}_T = \mathbf{c}_{\mathbf{x}_T} + \mathbf{C}_{\mathbf{x}_T, \mathbf{u}_T} \mathbf{k}_T + \mathbf{K}_T^T \mathbf{c}_{\mathbf{u}_T} + \mathbf{K}_T^T \mathbf{C}_{\mathbf{u}_T, \mathbf{u}_T} \mathbf{k}_T$$

# KMM: Start with Linear Dynamics

Solve for  $\mathbf{u}_{T-1}$  in terms of  $\mathbf{x}_{T-1}$

$\mathbf{u}_{T-1}$  affects  $\mathbf{x}_T$ !

$$f(\mathbf{x}_{T-1}, \mathbf{u}_{T-1}) = \mathbf{x}_T = \mathbf{F}_{T-1} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix} + \mathbf{f}_{T-1}$$

$$Q(\mathbf{x}_{T-1}, \mathbf{u}_{T-1}) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \mathbf{C}_{T-1} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix} + \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \mathbf{c}_{T-1} + \underbrace{V(f(\mathbf{x}_{T-1}, \mathbf{u}_{T-1}))}_{V(\mathbf{x}_T) = \text{const} + \frac{1}{2} \mathbf{x}_T^T \mathbf{V}_T \mathbf{x}_T + \mathbf{x}_T^T \mathbf{v}_T}$$

$$V(\mathbf{x}_T) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \underbrace{\mathbf{F}_{T-1}^T \mathbf{V}_T \mathbf{F}_{T-1}}_{\text{quadratic}} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix} + \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \underbrace{\mathbf{F}_{T-1}^T \mathbf{V}_T \mathbf{f}_{T-1}}_{\text{linear}} + \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \underbrace{\mathbf{F}_{T-1}^T \mathbf{v}_T}_{\text{linear}}$$

# KMM: Start with Linear Dynamics

Solve for  $\mathbf{u}_{T-1}$  in terms of  $\mathbf{x}_{T-1}$

$\mathbf{u}_{T-1}$  affects  $\mathbf{x}_T$ !

$$f(\mathbf{x}_{T-1}, \mathbf{u}_{T-1}) = \mathbf{x}_T = \mathbf{F}_{T-1} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix} + \mathbf{f}_{T-1}$$

$$Q(\mathbf{x}_{T-1}, \mathbf{u}_{T-1}) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \mathbf{C}_{T-1} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix} + \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \mathbf{c}_{T-1} + \underbrace{V(f(\mathbf{x}_{T-1}, \mathbf{u}_{T-1}))}_{V(\mathbf{x}_T) = \text{const} + \frac{1}{2} \mathbf{x}_T^T \mathbf{V}_T \mathbf{x}_T + \mathbf{x}_T^T \mathbf{v}_T}$$

$$V(\mathbf{x}_T) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \underbrace{\mathbf{F}_{T-1}^T \mathbf{V}_T \mathbf{F}_{T-1}}_{\text{quadratic}} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix} + \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \underbrace{\mathbf{F}_{T-1}^T \mathbf{V}_T \mathbf{f}_{T-1}}_{\text{linear}} + \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \underbrace{\mathbf{F}_{T-1}^T \mathbf{v}_T}_{\text{linear}}$$

# KMM: Start with Linear Dynamics

$$Q(\mathbf{x}_{T-1}, \mathbf{u}_{T-1}) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \mathbf{C}_{T-1} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix} + \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \mathbf{c}_{T-1} + V(f(\mathbf{x}_{T-1}, \mathbf{u}_{T-1}))$$

$$V(\mathbf{x}_T) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \underbrace{\mathbf{F}_{T-1}^T \mathbf{V}_T \mathbf{F}_{T-1}}_{\text{quadratic}} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix} + \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \underbrace{\mathbf{F}_{T-1}^T \mathbf{V}_T \mathbf{f}_{T-1}}_{\text{linear}} + \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \underbrace{\mathbf{F}_{T-1}^T \mathbf{v}_T}_{\text{linear}}$$

$$Q(\mathbf{x}_{T-1}, \mathbf{u}_{T-1}) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \mathbf{Q}_{T-1} \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix} + \begin{bmatrix} \mathbf{x}_{T-1} \\ \mathbf{u}_{T-1} \end{bmatrix}^T \mathbf{q}_{T-1}$$

$$\mathbf{Q}_{T-1} = \mathbf{C}_{T-1} + \mathbf{F}_{T-1}^T \mathbf{V}_T \mathbf{F}_{T-1}$$

$$\mathbf{q}_{T-1} = \mathbf{c}_{T-1} + \mathbf{F}_{T-1}^T \mathbf{V}_T \mathbf{f}_{T-1} + \mathbf{F}_{T-1}^T \mathbf{v}_T$$

$$\nabla_{\mathbf{u}_{T-1}} Q(\mathbf{x}_{T-1}, \mathbf{u}_{T-1}) = \mathbf{Q}_{\mathbf{u}_{T-1}, \mathbf{x}_{T-1}} \mathbf{x}_{T-1} + \mathbf{Q}_{\mathbf{u}_{T-1}, \mathbf{u}_{T-1}} \mathbf{u}_{T-1} + \mathbf{q}_{\mathbf{u}_{T-1}}^T = 0$$

$$\mathbf{u}_{T-1} = \mathbf{K}_{T-1} \mathbf{x}_{T-1} + \mathbf{k}_{T-1}$$

$$\mathbf{K}_{T-1} = -\mathbf{Q}_{\mathbf{u}_{T-1}, \mathbf{u}_{T-1}}^{-1} \mathbf{Q}_{\mathbf{u}_{T-1}, \mathbf{x}_{T-1}}$$

$$\mathbf{k}_{T-1} = -\mathbf{Q}_{\mathbf{u}_{T-1}, \mathbf{u}_{T-1}}^{-1} \mathbf{q}_{\mathbf{u}_{T-1}}$$

# KMM: Start with Linear Dynamics

Backward recursion

for  $t = T$  to 1:

$$\mathbf{Q}_t = \mathbf{C}_t + \mathbf{F}_t^T \mathbf{V}_{t+1} \mathbf{F}_t$$

$$\mathbf{q}_t = \mathbf{c}_t + \mathbf{F}_t^T \mathbf{V}_{t+1} \mathbf{f}_t + \mathbf{F}_t^T \mathbf{v}_{t+1}$$

$$Q(\mathbf{x}_t, \mathbf{u}_t) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{Q}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{q}_t$$

$$\mathbf{u}_t \leftarrow \arg \min_{\mathbf{u}_t} Q(\mathbf{x}_t, \mathbf{u}_t) = \mathbf{K}_t \mathbf{x}_t + \mathbf{k}_t$$


$$\mathbf{K}_t = -\mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t}^{-1} \mathbf{Q}_{\mathbf{u}_t, \mathbf{x}_t}$$

$$\mathbf{k}_t = -\mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t}^{-1} \mathbf{q}_{\mathbf{u}_t}$$

$$\mathbf{V}_t = \mathbf{Q}_{\mathbf{x}_t, \mathbf{x}_t} + \mathbf{Q}_{\mathbf{x}_t, \mathbf{u}_t} \mathbf{K}_t + \mathbf{K}_t^T \mathbf{Q}_{\mathbf{u}_t, \mathbf{x}_t} + \mathbf{K}_t^T \mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t} \mathbf{K}_t$$

$$\mathbf{v}_t = \mathbf{q}_{\mathbf{x}_t} + \mathbf{Q}_{\mathbf{x}_t, \mathbf{u}_t} \mathbf{k}_t + \mathbf{K}_t^T \mathbf{q}_{\mathbf{u}_t} + \mathbf{K}_t^T \mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t} \mathbf{k}_t$$

$$V(\mathbf{x}_t) = \text{const} + \frac{1}{2} \mathbf{x}_t^T \mathbf{V}_t \mathbf{x}_t + \mathbf{x}_t^T \mathbf{v}_t$$



Forward recursion

for  $t = 1$  to  $T$ :

$$\mathbf{u}_t = \mathbf{K}_t \mathbf{x}_t + \mathbf{k}_t$$

$$\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{u}_t)$$

# Model-Based: Further Reading

- PILCO: <http://mlg.eng.cam.ac.uk/pub/pdf/DeiRas11.pdf>
- DeepPilco: <http://mlg.eng.cam.ac.uk/yarin/PDFs/DeepPILCO.pdf>
- SOLAR: <https://arxiv.org/pdf/1808.09105.pdf>
- Gaussian Processes for RL:  
<https://papers.nips.cc/paper/2420-gaussian-processes-in-reinforcement-learning.pdf>
- LQR Sample Complexity: <https://arxiv.org/abs/1710.01688>
- Multi-step model-based: <https://arxiv.org/pdf/1811.00128.pdf>
- Theoretical guarantees for model-based RL: <https://openreview.net/pdf?id=BJe1E2R5KX>
- Guided Policy Search: [https://graphics.stanford.edu/projects/gpspaper/gps\\_full.pdf](https://graphics.stanford.edu/projects/gpspaper/gps_full.pdf)

# Approximate Dynamic Programming

- Directly approximates optimal control costs and solves this using techniques from dynamic programming

$$Q(x, u) = \max \left\{ \mathbb{E}_{e_t} \left[ \sum_{t=0}^N R(x_t, u_t) \right] : x_{t+1} = f(x_t, u_t, e_t), (x_0, u_0) = (x, u) \right\}$$

- Optimal policy therefore is

$$u_k = \arg \max_u Q_k(x_k, u)$$



# Approximate Dynamic Programming

$$Q_{\gamma}(x_k, u_k) \approx R(x_k, u_k) + \gamma \max_{u'} Q_{\gamma}(x_{k+1}, u').$$

With Update Rule


$$Q_{\gamma}^{(\text{new})}(x_k, u_k) = (1 - \eta)Q_{\gamma}^{(\text{old})}(x_k, u_k) + \eta \left( R(x_k, u_k) + \gamma \max_{u'} Q_{\gamma}^{(\text{old})}(x_{k+1}, u') \right)$$

Or alternatively, in the value iteration case

$$V(x) = \max_u Q(x, u).$$

# Approximate Dynamic Programming

online Q iteration algorithm:

- 
1. take some action  $\mathbf{a}_i$  and observe  $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$
  2.  $\mathbf{y}_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'} Q_\phi(\mathbf{s}'_i, \mathbf{a}'_i)$
  3.  $\phi \leftarrow \phi - \alpha \frac{dQ_\phi}{d\phi}(\mathbf{s}_i, \mathbf{a}_i)(Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{y}_i)$
- these are correlated!
- isn't this just gradient descent? that converges, right?

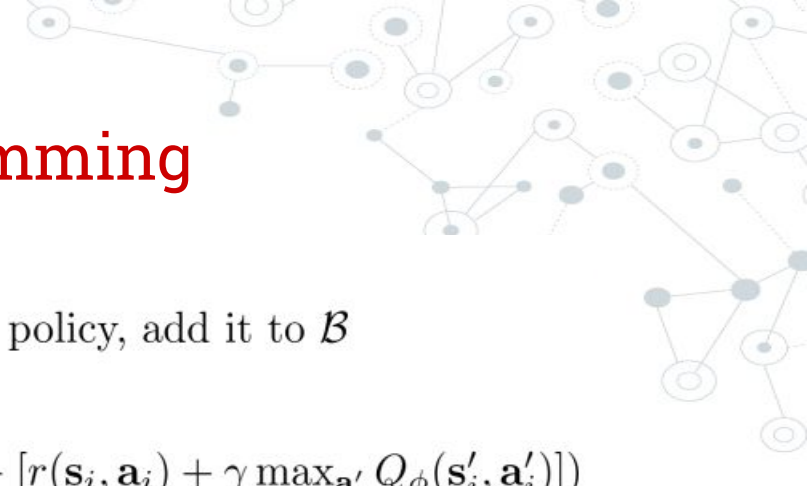

Q-learning is *not* gradient descent!

$$\phi \leftarrow \phi - \alpha \frac{dQ_\phi}{d\phi}(\mathbf{s}_i, \mathbf{a}_i)(Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - \underbrace{(r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'} Q_\phi(\mathbf{s}'_i, \mathbf{a}'_i))}_{\text{no gradient through target value}})$$

no gradient through target value

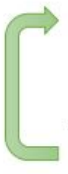

# Approximate Dynamic Programming

full Q-learning with replay buffer:

- 
- 
1. collect dataset  $\{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)\}$  using some policy, add it to  $\mathcal{B}$
  2. sample a batch  $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$  from  $\mathcal{B}$
  3.  $\phi \leftarrow \phi - \alpha \sum_i \frac{dQ_\phi}{d\phi}(\mathbf{s}_i, \mathbf{a}_i)(Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - [r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'} Q_\phi(\mathbf{s}'_i, \mathbf{a}'_i)])$

Stabilize regression with

full fitted Q-iteration algorithm:

- 
- 
1. collect dataset  $\{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)\}$  using some policy
  2. set  $\mathbf{y}_i \leftarrow r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'_i} Q_\phi(\mathbf{s}'_i, \mathbf{a}'_i)$
  3. set  $\phi \leftarrow \arg \min_{\phi} \frac{1}{2} \sum_i \|Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{y}_i\|^2$

**perfectly well-defined, stable regression**

# DQN (Mnih et. al 2015)

Q-learning with replay buffer and target network:

1. save target network parameters:  $\phi' \leftarrow \phi$
2. collect dataset  $\{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)\}$  using some policy, add it to  $\mathcal{B}$
- $N \times$   
 $K \times$  3. sample a batch  $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$  from  $\mathcal{B}$
4.  $\phi \leftarrow \phi - \alpha \sum_i \frac{dQ_\phi}{d\phi}(\mathbf{s}_i, \mathbf{a}_i)(Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - [r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'} Q_{\phi'}(\mathbf{s}'_i, \mathbf{a}')])$



## Further Reading

- DQN: <https://www.cs.toronto.edu/~vmnih/docs/dqn.pdf>
- Double DQN: <https://arxiv.org/abs/1509.06461>
- Dueling DQN: <https://arxiv.org/abs/1511.06581>
- DDPG: <https://arxiv.org/pdf/1509.02971.pdf>
- Is Q-learning Provably Efficient: <https://arxiv.org/pdf/1807.03765.pdf>

# Direct Policy Search

- The idea is, use as little information as possible, still find a policy.

$$\text{maximize}_{\vartheta} \quad \mathbb{E}_{p(z;\vartheta)}[R(z)] .$$

- So we try to maximize a cost function

$$J(\vartheta) := \mathbb{E}_{p(z;\vartheta)}[R(z)]$$

- And with a touch of nice math, we get

$$\begin{aligned} \nabla_{\vartheta} J(\vartheta) &= \int R(z) \nabla_{\vartheta} p(z; \vartheta) dz \\ &= \int R(z) \left( \frac{\nabla_{\vartheta} p(z; \vartheta)}{p(z; \vartheta)} \right) p(z; \vartheta) dz \\ &= \int (R(z) \nabla_{\vartheta} \log p(z; \vartheta)) p(z; \vartheta) dz \\ &= \mathbb{E}_{p(z;\vartheta)} [R(z) \nabla_{\vartheta} \log p(z; \vartheta)] . \end{aligned}$$

# Direct Policy Search

- The idea is, use as little information as possible, still find a policy.

$$\text{maximize}_{\vartheta} \quad \mathbb{E}_{p(z;\vartheta)}[R(z)] .$$

- So we try to maximize a cost function

$$J(\vartheta) := \mathbb{E}_{p(z;\vartheta)}[R(z)]$$

- And with a touch of nice math, we get

$$\begin{aligned} \nabla_{\vartheta} J(\vartheta) &= \int R(z) \nabla_{\vartheta} p(z; \vartheta) dz \\ &= \int R(z) \left( \frac{\nabla_{\vartheta} p(z; \vartheta)}{p(z; \vartheta)} \right) p(z; \vartheta) dz \\ &= \int (R(z) \nabla_{\vartheta} \log p(z; \vartheta)) p(z; \vartheta) dz \\ &= \mathbb{E}_{p(z;\vartheta)} [R(z) \nabla_{\vartheta} \log p(z; \vartheta)] . \end{aligned}$$

# Direct Policy Search

---

## Algorithm 1 REINFORCE

---

- 1: **Hyperparameters:** step-sizes  $\alpha_j > 0$ .
  - 2: **Initialize:**  $\vartheta_0$  and  $k = 0$ .
  - 3: **while** ending condition not satisfied **do**
  - 4:   Sample  $z_k \sim p(z; \vartheta_k)$ .
  - 5:   Set  $\vartheta_{k+1} = \vartheta_k + \alpha_k R(z_k) \nabla_{\vartheta} \log p(z_k; \vartheta_k)$ .
  - 6:    $k \leftarrow k + 1$
  - 7: **end while**
-



# Pure Random Search

- All we do is perturb decision variable  $z$ , and see what happens
- Sample from some distribution

$$p(z; \vartheta) = p_0(z - \vartheta)$$

- With a gradient estimation

$$g_{\sigma}(\vartheta) = \frac{R(\vartheta + \sigma\epsilon) - R(\vartheta - \sigma\epsilon)}{2\sigma} \epsilon.$$

- And averaging

$$g_{\sigma}^{(m)}(\vartheta) = \frac{1}{m} \sum_{i=1}^m \frac{R(\vartheta + \sigma\epsilon_i) - R(\vartheta - \sigma\epsilon_i)}{2\sigma} \epsilon_i.$$

# Policy Gradients: Further Reading

- Random Search is competitive in RL: <https://arxiv.org/abs/1803.07055>
- TRPO: <https://arxiv.org/abs/1502.05477>
- PPO: <https://arxiv.org/pdf/1707.06347.pdf>

# Further Reading

- RL Book: <http://incompleteideas.net/book/bookdraft2017nov5.pdf>
- Neuro-dynamic programming: <http://athenasc.com/ndpbook.html>

# Thanks!

**Congratulations, you've made  
it to the end :)**

